

Performance Tuning for Apache Tomcat

Mark Thomas
April 2009

Who am I?

Apache Tomcat committer

Resolved 1,500+ Tomcat bugs

Apache Tomcat PMC member

Member of the Apache Software Foundation

Member of the ASF security committee

Created the Tomcat security pages

Senior Software Engineer and Consultant at
SpringSource

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache

- JVM

Scaling Tomcat

Hints and tips

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache

- JVM

Scaling Tomcat

Hints and tips

The process

Understand the system architecture

Stabilise the system

Set the performance target(s)

Measure current performance

Identify the current bottleneck

Fix the root cause of the bottleneck

Repeat until you meet the target

Common errors

Optimising code that doesn't need it

Insufficient testing

- realistic data volumes

- realistic user load

Lack of clear performance targets

Guessing where the bottleneck is

Fixing the symptom rather than the cause

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache

- JVM

Scaling Tomcat

Hints and tips

Tomcat tuning

Applications typically account for >80% of request processing time

Remember the tuning process

Focus your efforts on the bottlenecks

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging**

- connectors

- content cache

- JVM

Scaling Tomcat

Hints and tips

Production logging

Default configuration is generic

Some settings not ideal for production

- catch-all logger logs to file and stdout
- no overflow protection

Production logging

Remove duplicate logging (logging.properties)

```
.handlers = 1catalina.org.apache.juli.FileHandler,  
            java.util.logging.ConsoleHandler
```

becomes

```
.handlers = 1catalina.org.apache.juli.FileHandler
```

To add rotation

```
1catalina.java.util.logging.FileHandler.pattern =  
    ${catalina.base}/logs/catalina.%g.log  
1catalina.java.util.logging.FileHandler.limit =  
20000000  
1catalina.java.util.logging.FileHandler.count = 5
```

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors**

- content cache

- JVM

Scaling Tomcat

Hints and tips

Connector tuning

Depends on

- your application usage patterns

- your network

- TCP connections

- HTTP transactions

- HTTP Keep-Alive

- SSL

Additional considerations for load balancing

- Layer 4 or Layer 7

- Connection pools

Connector tuning

HTTP/1.0 has no keep alive

- Client creates TCP connection to server

- Client sends request

- Server sends response

- Connection closed

Modern web pages can require >100 requests to completely display the page

Creating TCP connections can be expensive

- Unlikely to be an issue on a LAN

- May well be an issue for mobile devices

Connector tuning

HTTP/1.1 introduced keep alive

- Client creates TCP connection to server

- Client sends first request

- Server sends first response

- Client sends second request

- Server sends second response

- ...

- Connection closed

Connectors that use blocking IO use a thread to maintain the keep alive connection

Connector tuning

Layer 4 load balancer

- Does not understand HTTP

- Makes decisions based on client IP and port

Layer 7 load balancer

- Understands HTTP

- Can use HTTP headers to make decisions

Which connector?

Tomcat has a choice of three

Java Blocking IO

- Oldest – most stable

- JSSE based SSL

Native (APR)

- Non-blocking

- Uses OpenSSL

Java Non-blocking IO

- JSSE based SSL

Which connector?

Requirement	Connectors in preference order		
Stability	BIO	APR/NIO	
SSL	APR	NIO	BIO
Low concurrency	BIO	APR	NIO
High concurrency No Keep-Alive	BIO	APR	NIO
High concurrency Keep-Alive	APR	NIO	BIO

Which connector?

Why would you use the NIO connector?

The Native (APR) connector is unstable on Solaris

NIO is a pure Java solution

It is easy to switch between NIO and BIO with SSL

Connector tuning

maxThreads

maximum number of concurrent requests

for BIO, maximum number of open/active connections

typical values 200 to 800

400 is a good starting value

heavy CPU usage → decrease

light CPU usage → increase

Connector tuning

maxKeepAliveRequests

typical values 1 or 100

maximum number of HTTP requests per TCP connection

set to 1 to disable keep alive

disable for BIO with very high concurrency, layer 4 load balancer, no SSL

enable for SSL, APR/NIO, layer 7 load balancer

Note BIO connector automatically disables keep alive when concurrent connections reach 75% of maxThreads

Connector tuning

connectionTimeout

- typical value 3000

- default of 20000 is too high for production use

- also used for keep alive time-out

- increase for slow clients

- increase for layer 7 load balancer with connection pool
and keep alive on

- decrease for faster time-outs

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache**

- JVM

Scaling Tomcat

Hints and tips

Content cache tuning

Dynamic content is not cached

Static content is cached

Configured using the `<Context .../>` element

CacheMaxSize (KB)

10240

CacheTTL (ms)

5000

CacheMaxFileSize (KB) (6.0.19 onwards)

512

NIO/APR can use `SEND_FILE`

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache

- JVM**

Scaling Tomcat

Hints and tips

JVM tuning

Two key areas

- Memory

- Garbage collection

They are related

Remember to follow the tuning process

JVM tuning

Java heap (Xmx, Xms) is not the same as the process heap

Process heap includes

- Java Heap

- Permanent Generation

- Thread stacks

- Native code

- Directly allocated memory

- Code generation

- Garbage collection

- TCP buffers

Read OutOfMemory exception messages carefully

JVM tuning: memory

-Xms/-Xmx

Used to define size of Java heap

Aim to set as low as possible

Setting too high can cause wasted memory and long GC cycles

-XX:NewSize/-XX:NewRatio

Set to 25-33% of total Java heap

Setting too high or too low leads to inefficient GC

JVM tuning: ideal garbage collection

Short lived objects never reach the Old Generation

Short lived objects cleaned up by short minor
garbage collections

Long lived objects promoted to Old Generation

Long lived objects cleaned up by (rare) full garbage
collection

JVM tuning: garbage collection

GC pauses the application

Regardless of GC algorithm

Pause can range from milliseconds to seconds

The pause will impact your response time

How much does this matter?

`-XX:MaxGCPauseMillis` `-XX:MaxGCMinorPauseMillis`

Set GC pause time goals

More frequent GC, shorter pauses

JVM tuning: garbage collection

There are many more options

Useful reference

<http://blogs.sun.com/watt/resource/jvm-options-list.html>

Newer GC algorithms may not behave the way you expect

Concurrent Mark Sweep

- XX:+UseConcMarkSweepGC

- Does not use survivor spaces

- Can be forced to; not recommended

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache

- JVM

Scaling Tomcat

Hints and tips

Scaling Tomcat

Load balancing

Routing requests to multiple Tomcat instances

Clustering

Sharing state between Tomcat instances for fail-over

Scaling Tomcat

Simplest configuration

- 1 * httpd
- 2 * Tomcat instances
- mod_proxy_http

Considerations

- state management
- fail over

Scaling Tomcat

Stateless load-balancing using httpd

In httpd.conf:

```
# Cluster definition
<Proxy balancer://devcluster>
    BalancerMember http://192.168.0.31:8080 disablereuse=On
    BalancerMember http://192.168.0.32:8080 disablereuse=On
</Proxy>
# Route all requests to the cluster
ProxyPass / balancer://devcluster/
```

Scaling Tomcat

Enabling the manager In httpd.conf

```
# Pass all requests except the manager to the cluster
ProxyPass /balancer-manager !
ProxyPass / balancer://devcluster/
# Configure the manager
<Location /balancer-manager>
    SetHandler balancer-manager
    Order Deny,Allow
    Deny from all
    Allow from 127.0.0.1
</Location>
```

Scaling Tomcat

Add sticky session support
Tomcat configuration
server.xml

```
<Engine jvmRoute="tc01"... />
```

jvmRoute must be unique for each instance
httpd configuration

Add route to each balancer member

```
BalancerMember http://192.168.0.31:8080 disablereuse=On route=tc01
```

Configure sticky sessions on ProxyPass

```
ProxyPass / balancer://devcluster/  
nofailover=On stickysession=JSESSIONID|jsessionid
```

Scaling Tomcat

Add session replication

Application configuration (WEB-INF/web.xml)

- Add `<distributable/>`

- Keep the session as small as possible

- Session attributes must implement Serializable

Tomcat configuration (server.xml)

- Uncomment `<Cluster ... />` element under `<Engine ... >`

- Defaults to get you started

- Overview: </docs/cluster-howto.html>

- Details: </docs/config/cluster.html>

Scaling Tomcat

httpd configuration (httpd.conf)

```
ProxyPass / balancer://devcluster/  
           nofailover=Off stickysession=JSESSIONID|jsessionid
```

Fail over

- Session replication asynchronous by default so usually used with sticky sessions

- Single line configuration for defaults

- Uses multicast for node discovery

- Will need additional configuration for production use

Agenda

The optimisation / tuning process

Tomcat tuning options

- logging

- connectors

- content cache

- JVM

Scaling Tomcat

Hints and tips

Hints and tips

Load balancing / clustering

- use a minimum of 3 Tomcat instances

- use load balancing and clustering in your development environment

Redeployment can expose memory leaks

- include this in your testing

Remember to follow the process

Questions?

mark.thomas@springsource.com

<http://www.springsource.com/webinars>

markt@apache.org

users@tomcat.apache.org