

Tomcat Expert Series

Performance Tuning

Filip Hanik
SpringSource
2009

- Performance Tuning Process
- Logging improvements
- TCP and HTTP
- Tuning your connectors
- Content delivery and caching
- Tuning the JVM

- Understand the system architecture
- Stabilise the system
- Set Performance Targets
 - Web applications are easy
 - Only one consideration – request/response time

- Measure current performance
- Identify the current bottleneck
 - Focus on one item at a time
- Fix the **root** cause
 - Easy to get side tracked

- When possible, tune pre-production
 - Hard to profile in production
- Application tuning most important
 - 80% or more of request time is typically spent inside the application
- Tomcat tuning is fairly limited
 - Divided between JVM tuning and Tomcat connectors
 - Requires lower level of understanding

- Out of the Box Tomcat
 - Tomcat is ready for production
- JVM settings must be applied
 - Default memory settings usually too small for most web applications
- Tuning is limited
 - So we can cover most of it

- Tomcat logging is fairly good
 - Years of adjusting log levels pays off
 - Doesn't log what you don't need to see
- A few gotcha's with the default configuration
 - Catch all logger creates duplicate logs
 - Standard out – often piped to catalina.out
 - Log file on the file system
 - Synchronized logging
 - No overflow protection

- Tomcat's logger
 - Rotated based on date
 - Implements a per-class-loader logger
 - Simply drop logging.properties into your web application and logging is configured
 - Synchronous logging
 - No file limit

- Java Virtual Machine logger
 - Rotated based on size
 - One global configuration for entire JVM
 - Synchronous logging

Remove duplicate logging (logging.properties)

```
.handlers = 1catalina.org.apache.juli.FileHandler,  
            java.util.logging.ConsoleHandler
```

Adjusted catch all logger

```
.handlers = 1catalina.org.apache.juli.FileHandler
```

- Overflow protection
 - Size based rotation using JVM logger

```
handlers = 1catalina.java.util.logging.FileHandler,...
```

- No more than 5x20mb files

```
1catalina.java.util.logging.FileHandler.pattern =  
    ${catalina.base}/logs/catalina.%g.log  
1catalina.java.util.logging.FileHandler.limit = 20000000  
1catalina.java.util.logging.FileHandler.count = 5
```

- Tuning Tomcat connectors
 - server.xml
 - <Connector>
- To properly tune one must
 - Understand the TCP protocol
 - Understand how the CPU works
 - Understand load balancers and their algorithms

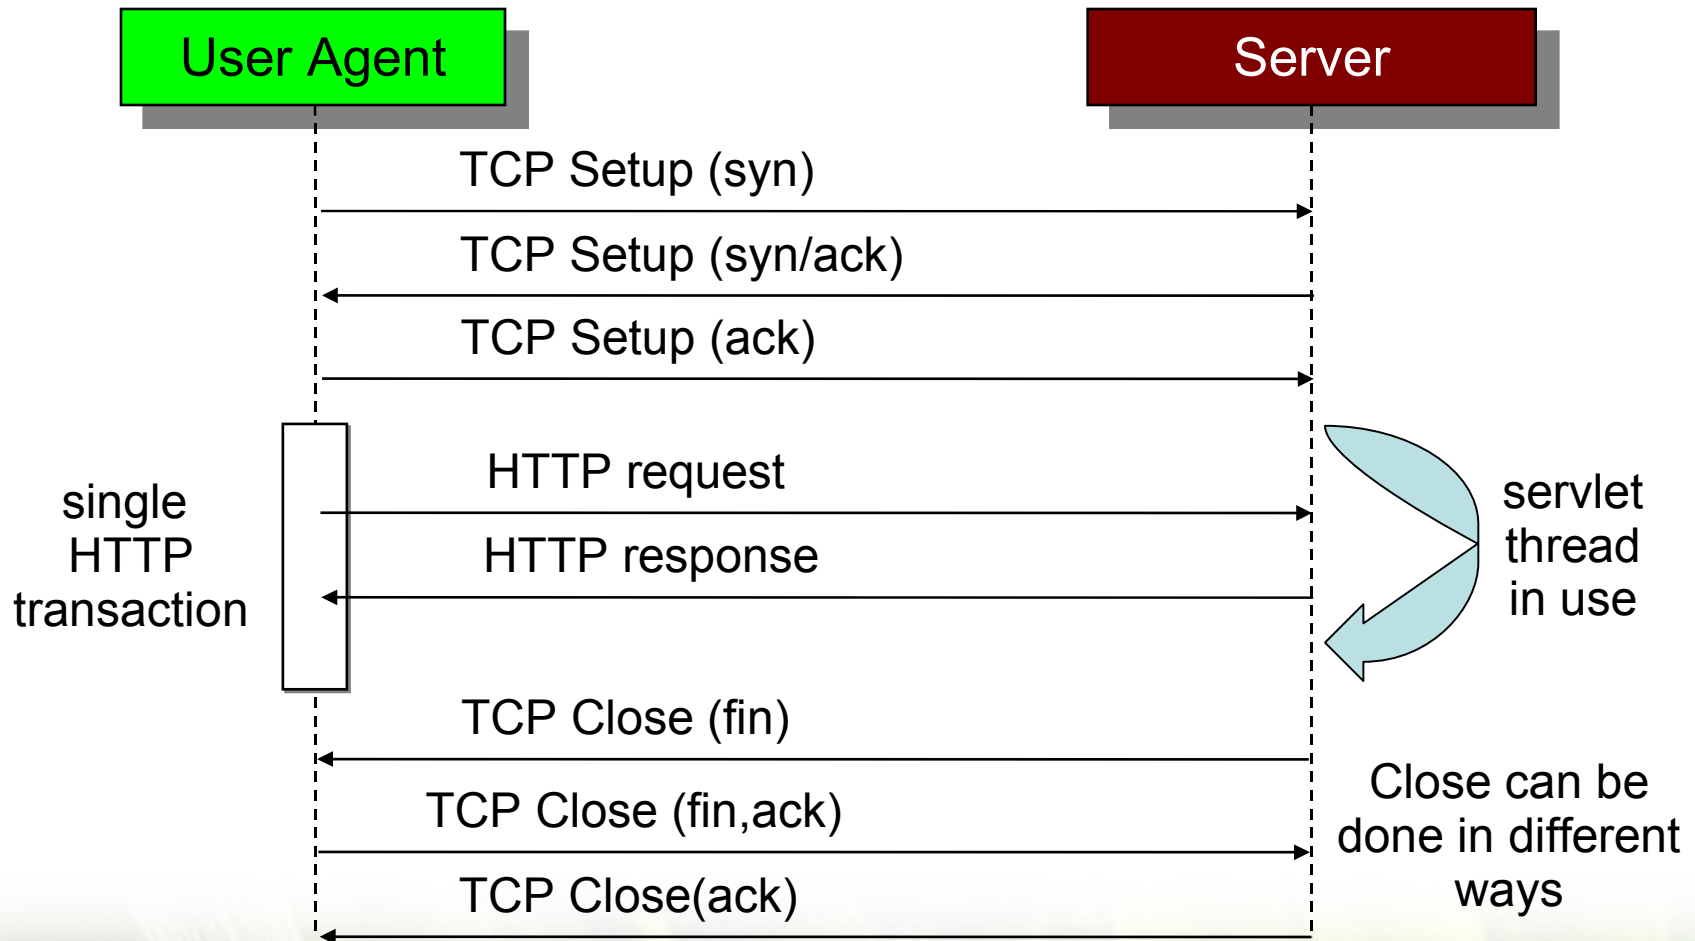
- Layer 4 in the OSI stack
- Session based
- TCP stack implementation keeps state
- Flow control
- Delivery guarantee

- Setup and break down handshakes
- Client response time
 - Handshakes add to HTTP transaction time
 - HTTP keep alive improves client response time
 - HTTP keep alive takes up server resources

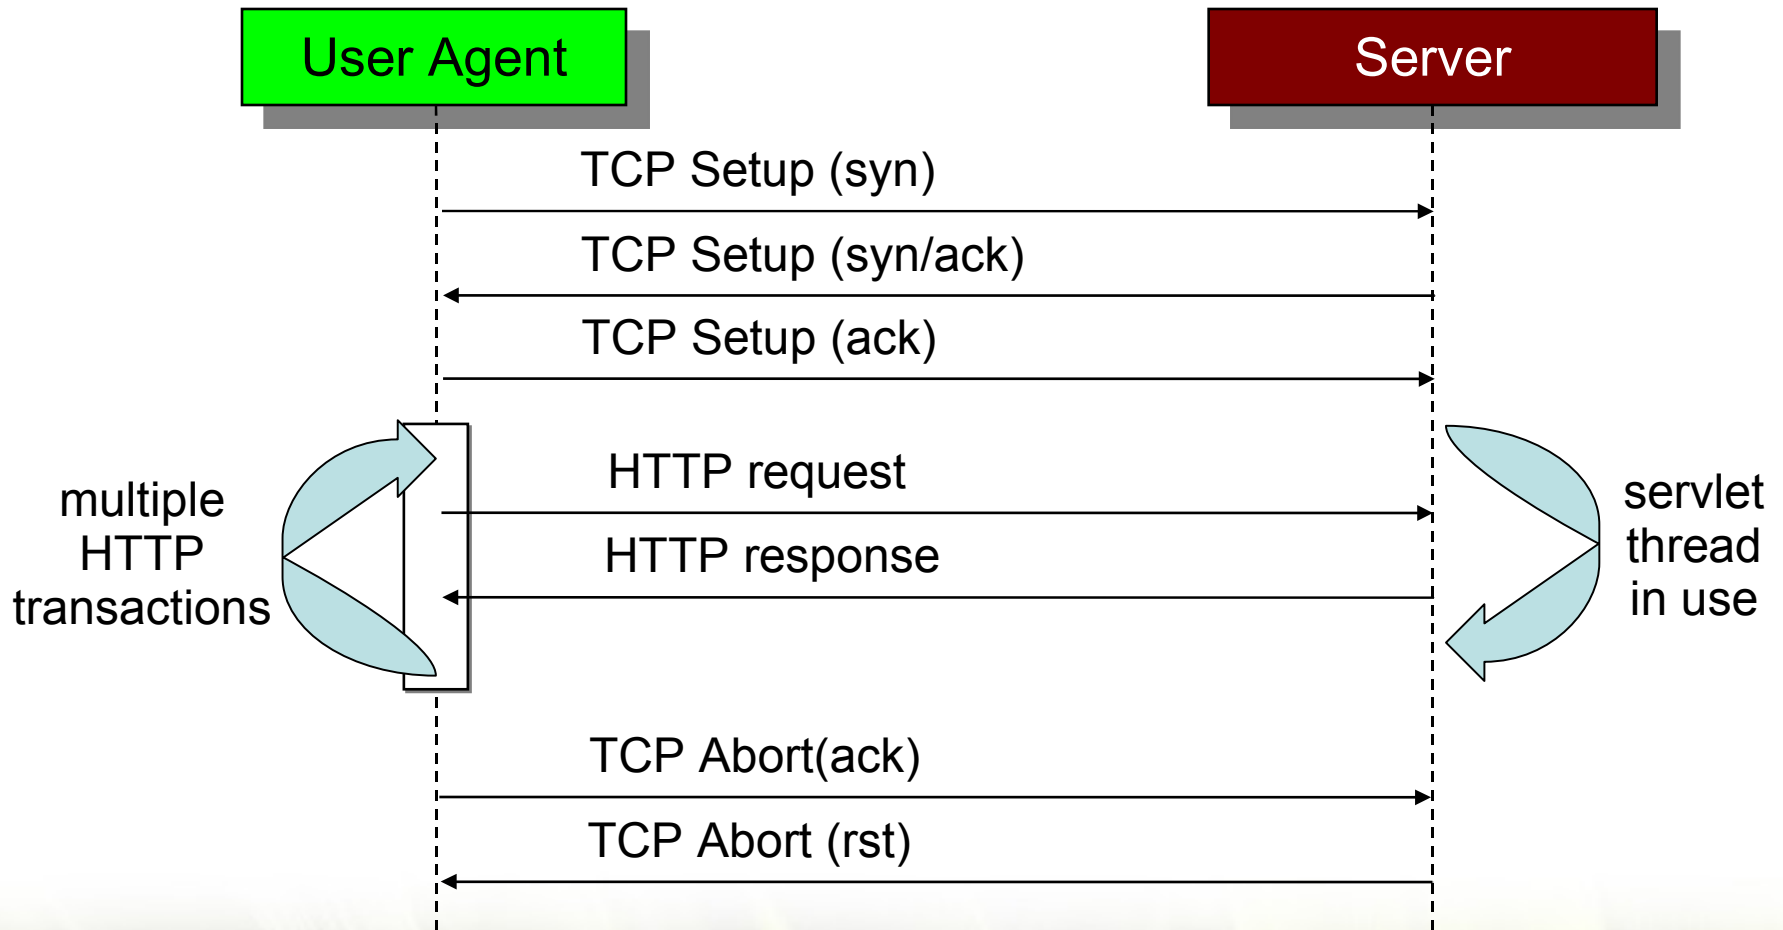
- Each connection represented by
 - source address
 - source port
 - destination address
 - destination port
- This is all the information a layer 4 load balancer has for load balancing

- Prevents buffer overflow and lost data
- Server must adjust write speed to client's ability to read
- Servlet specification is blocking IO
 - Utilize a thread for the entire HTTP transaction
 - For static content, Tomcat offers SEND_FILE with APR and NIO connectors

TCP: No Keep-Alive



TCP: Keep-Alive



- How does TCP affect our system?
 - Traffic patterns
 - High concurrency/short requests
 - Low concurrency/long requests
 - Static content
 - Dynamic content
 - Combination
 - Variations
 - Average size of request
 - It's these patterns that decide how to tune our system

-
- Layer 7 in the OSI stack
 - Stateless protocol

- HTTP over SSL
- Expensive handshake
 - Keep alive makes a big difference
- Encryption hides HTTP from routing devices
- For any appliances, such as LB, this means
 - Fallback to layer 4

- TCP
 - Based on destination address/port
 - Connection centric – 1:1
 - Can lead to uneven loads
- HTTP
 - Based on HTTP headers
 - Can reuse server connections
 - Can drain sessions

- Load balancing
 - Connection limits
 - Reusing connections
 - Traffic shaping
- Load balancing algorithm drive Tomcat configuration choices

- Our tuning options
 - Threads
 - Keep alive requests
 - TCP Backlog (acceptCount)
 - connectionTimeout
 - Socket buffers
- Different connectors
 - BIO – Blocking Java connector, default
 - APR – Uses native C code for IO
 - NIO – Non blocking Java connectors

- Disclaimer
 - Tuning options are meant for working and high performing applications
 - Options will not fix bad application behavior
 - If application is not tuned
 - Situation can worsen

Which connector?

- Use BIO if:
 - Stability is the highest priority
 - APR and NIO are more recent
 - Most content is dynamic
 - Keep alive is not a determining factor

```
protocol="org.apache.coyote.http11.Http11Protocol"
```

Which connector?

- Use APR if:
 - SSL is terminated at Tomcat
 - Keep alive is important
 - Lots of static content
 - Using Comet feature
 - Requires compilation of native library

```
protocol="org.apache.coyote.http11.Http11AprProtocol"
```

Which connector?

- Use NIO if:
 - Compiling APR is not an option
 - Keep alive is important
 - Using SSL
 - Lots of static content
 - Using Comet features

```
protocol="org.apache.coyote.http11.Http11NioProtocol"
```

Which connector?

- If uncertain:
 - Use BIO connector
 - Most mature code, both in Tomcat and JVM
 - Will not break down
 - Auto tune feature to disable keep alive
 - When hitting 75% of maxThreads in connection count

```
protocol="org.apache.coyote.http11.Http11Protocol"
```

Which connector?

Comparison Chart	<u>Java BIO</u>	<u>Java NIO</u>	<u>APR</u>
<u>Class</u>	Http11Protocol	Http11NioProtocol	Http11AprProtocol
<u>Version</u>	3.x+	6.x+	5.5.x+
<u>Polling</u>	NO	YES	YES
<u>Polling Size</u>	N/A	Unlimited	Configurable
<u>HTTP Req Read</u>	Blocking	Non blocking	Blocking
<u>HTTP Body Read</u>	Blocking	Sim Blocking	Blocking
<u>HTTP Write</u>	Blocking	Sim Blocking	Blocking
<u>SSL</u>	JSSE	JSSE	OpenSSL
<u>SSL Handshake</u>	Blocking	Non blocking	Blocking
<u>Max Connections</u>	maxThreads	Unlimited	Configurable

Which connector?

- If uncertain:
 - Use BIO connector
 - Most mature code, both in Tomcat and JVM
 - Will not break down
 - Auto tune feature to disable keep alive
 - When hitting 75% of maxThreads in connection count

```
protocol="org.apache.coyote.http11.Http11Protocol"
```

Which connector?

- If uncertain:
 - Use BIO connector
 - Most mature code, both in Tomcat and JVM
 - Will not break down
 - Auto tune feature to disable keep alive
 - When hitting 75% of maxThreads in connection count

```
protocol="org.apache.coyote.http11.Http11Protocol"
```


- **maxThreads**
 - Typical range 200–800
 - Maximum nr of concurrent requests
 - For BIO, max nr of open/active connections
 - Good starting value 400

- `maxThreads="400"`
 - Decrease if you see heavy CPU usage
 - Application might be CPU bound instead of IO bound
 - Find out what is causing CPU usage
 - Increase if you don't see much CPU usage
 - Applications could be synchronized -> no gain
 - Take into account other resources, such as database connections

- `maxKeepAliveRequests`
 - Typical values 1, 100–200
 - Represents the number of requests Tomcat will handle on a TCP connection
 - Set to 1 disables keep alive
 - `connectionTimeout/keepAliveTimeout` controls the timeout in between requests

- maxKeepAliveRequests
 - Set to 1 if
 - Very high concurrency
 - Not using SSL in Tomcat
 - Using layer 4 load balancer
 - Using BIO connector
 - Set to >1 if
 - Using SSL or low concurrency
 - Layer 7 load balancer with advanced features
 - Using APR or NIO connector
 - BIO connector automatically disables keep alive for high connection counts

- `acceptCount`
 - Typical ranges 50–300
 - Represents nr of additional connections the OS should accept on your behalf
 - Useful when Tomcat can't accept connections fast enough

- `acceptCount="100"`
 - Increase if
 - Very high concurrency (nr of connections)
 - Connections getting rejected during peak traffic
 - Keep alive should be off
 - Decrease if
 - Keep alive is on
 - Connections getting accepted but never serviced

- `connectionTimeout`
 - Values from 2000–60000
 - Represents the `SO_TIMEOUT` value
 - Essentially, max time between TCP packets during a blocking read or write
 - Critical to a stable system
 - Also used for keep alive timeout

- `connectionTimeout="3000"`
 - Increase if
 - Working with slow clients (dial up)
 - Using a layer 7 load balancer with connection limit/pool and keep alive on
 - Decrease if
 - Need faster timeouts
 - Default value of 20,000 (20secs) is too high for a web server

- Dynamic content
 - No caching done
 - Tomcat has to deliver it blocking mode
 - Worker thread is not released until all content has been delivered
 - Fast dynamic content can piggy back on send file
 - Simply write to file, set request attribute and hand off to Tomcat's poller threads

- Static content
 - Size based cache, default 10mb
 - BIO – Tomcat has to deliver it blocking mode
 - NIO/APR
 - Tomcat can use SEND_FILE
 - Release worker thread, deliver the content using a background thread when the client is ready to receive

- Configured in `<Context>` element
- 40MB cache (default 10MB)
- cache revalidation every 60 seconds (default 5 seconds)
- caching enabled (default true)

```
<Context cacheMaxSize="40960" cacheTTL="60000"  
          cachingAllowed="true">  
</Context>
```

- Key parameters for JVM tuning
 - Memory
 - Garbage collection
- They are not independent

- Short lived objects never reach the Old Generation
- Short lived objects cleaned up by short minor garbage collections
- Long lived objects promoted to Old Generation
- Long lived objects cleaned up by (rare) full garbage collection

- `-Xms/-Xmx`
 - Used to define size of Java heap
 - Aim to set as low as possible
 - Setting too high can cause wasted memory and long GC cycles
- `-XX:NewSize/-XX:NewRatio`
 - Set to 25–33% of total Java heap
 - Setting too high or too low leads to inefficient GC
 - Often these are not tuned, GC log will reveal

- GC pauses the application
 - Regardless of GC algorithm
- Pause can range from milliseconds to seconds
- The pause will impact your response time
 - How much does this matter?

- –XX:MaxGCPauseMillis
 - Set GC pause time goal
 - More frequent GC
 - Shorter pauses
 - Goal is for major collections
- –XX:MaxGCMinorPauseMillis
 - Applies to young generation

- GC Settings – JDK 1.5 and 1.6

- XX:+UseConcMarkSweepGC
- XX:+CMSIncrementalMode
- XX:+CMSIncrementalPacing
- XX:CMSIncrementalDutyCycleMin=0
- XX:CMSIncrementalDutyCycle=10
- XX:+UseParNewGC
- XX:+CMSPermGenSweepingEnabled
- XX:+CMSClassUnloadingEnabled
- XX:MaxGCPauseMillis=250
- XX:MaxGCMinorPauseMillis=100

- Much bigger topic
- Same tuning rules apply
 - Doesn't compensate for bad, slow or poorly written applications
- Sun JDK options

<http://blogs.sun.com/watt/resource/jvm-options-list.html>

Questions...

and answers!