

Being Ready for Apache Kafka: Today's Ecosystem and Future Roadmap

Michael G. Noll

@miguno


Developer Evangelist, Confluent Inc.



- Developer Evangelist at Confluent since August '15
- Previously Big Data lead at .COM/.NET DNS operator Verisign
- Blogging at <http://www.michael-noll.com/> (too little time!)
- PMC member of Apache Storm (too little time!)
- michael@confluent.io



- Founded in Fall 2014 by the creators of Apache Kafka
- Headquartered in San Francisco bay area
- We provide a stream data platform based on Kafka
- We contribute **a lot** to Kafka, obviously 😊



kafka



Apache Kafka is the **distributed, durable equivalent** of Unix pipes.
Use it to connect and compose your large-scale data apps.

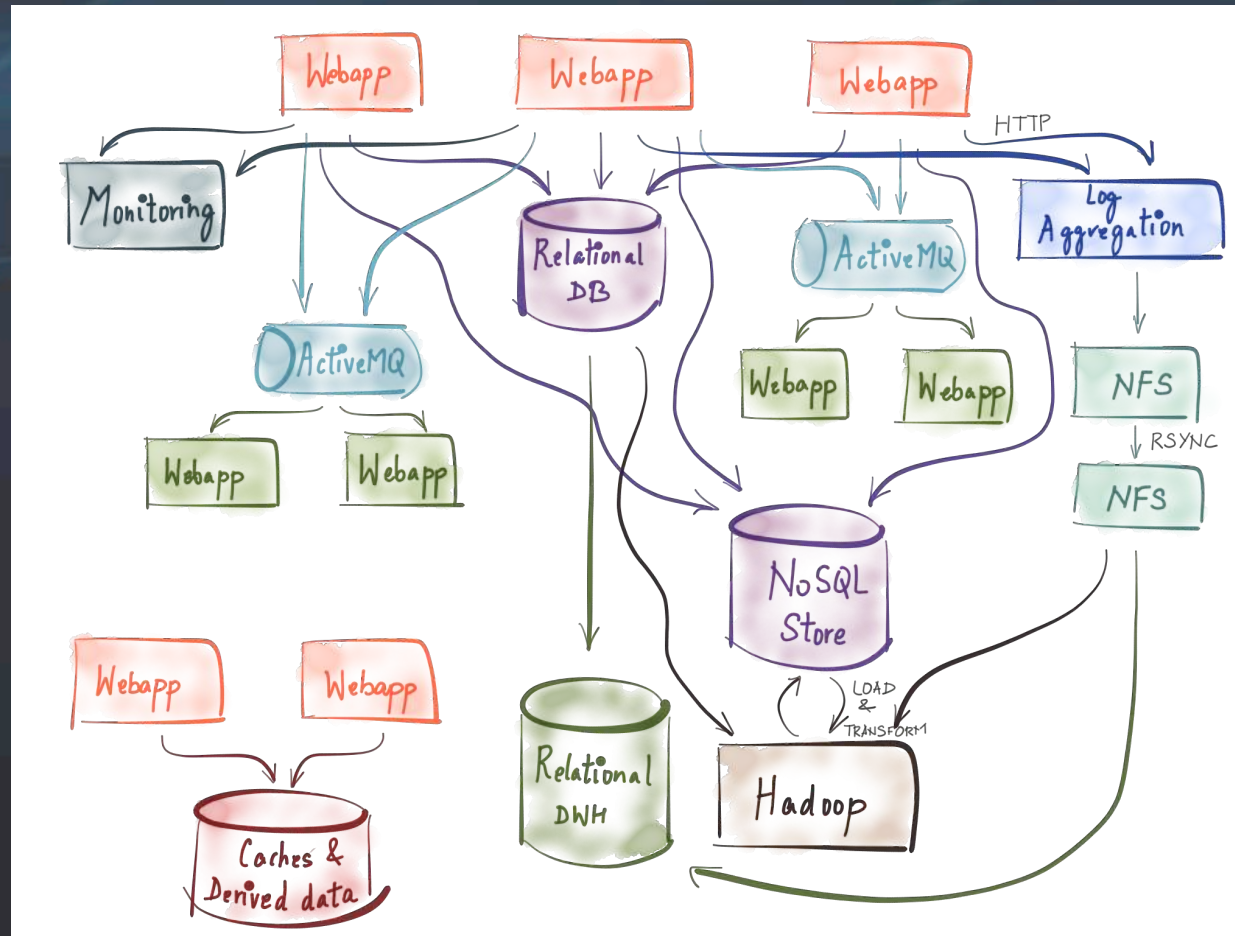
```
$ cat < in.txt | grep "apache" | tr a-z A-Z > out.txt
```




this

this

Example: LinkedIn before Kafka



Example: LinkedIn after Kafka

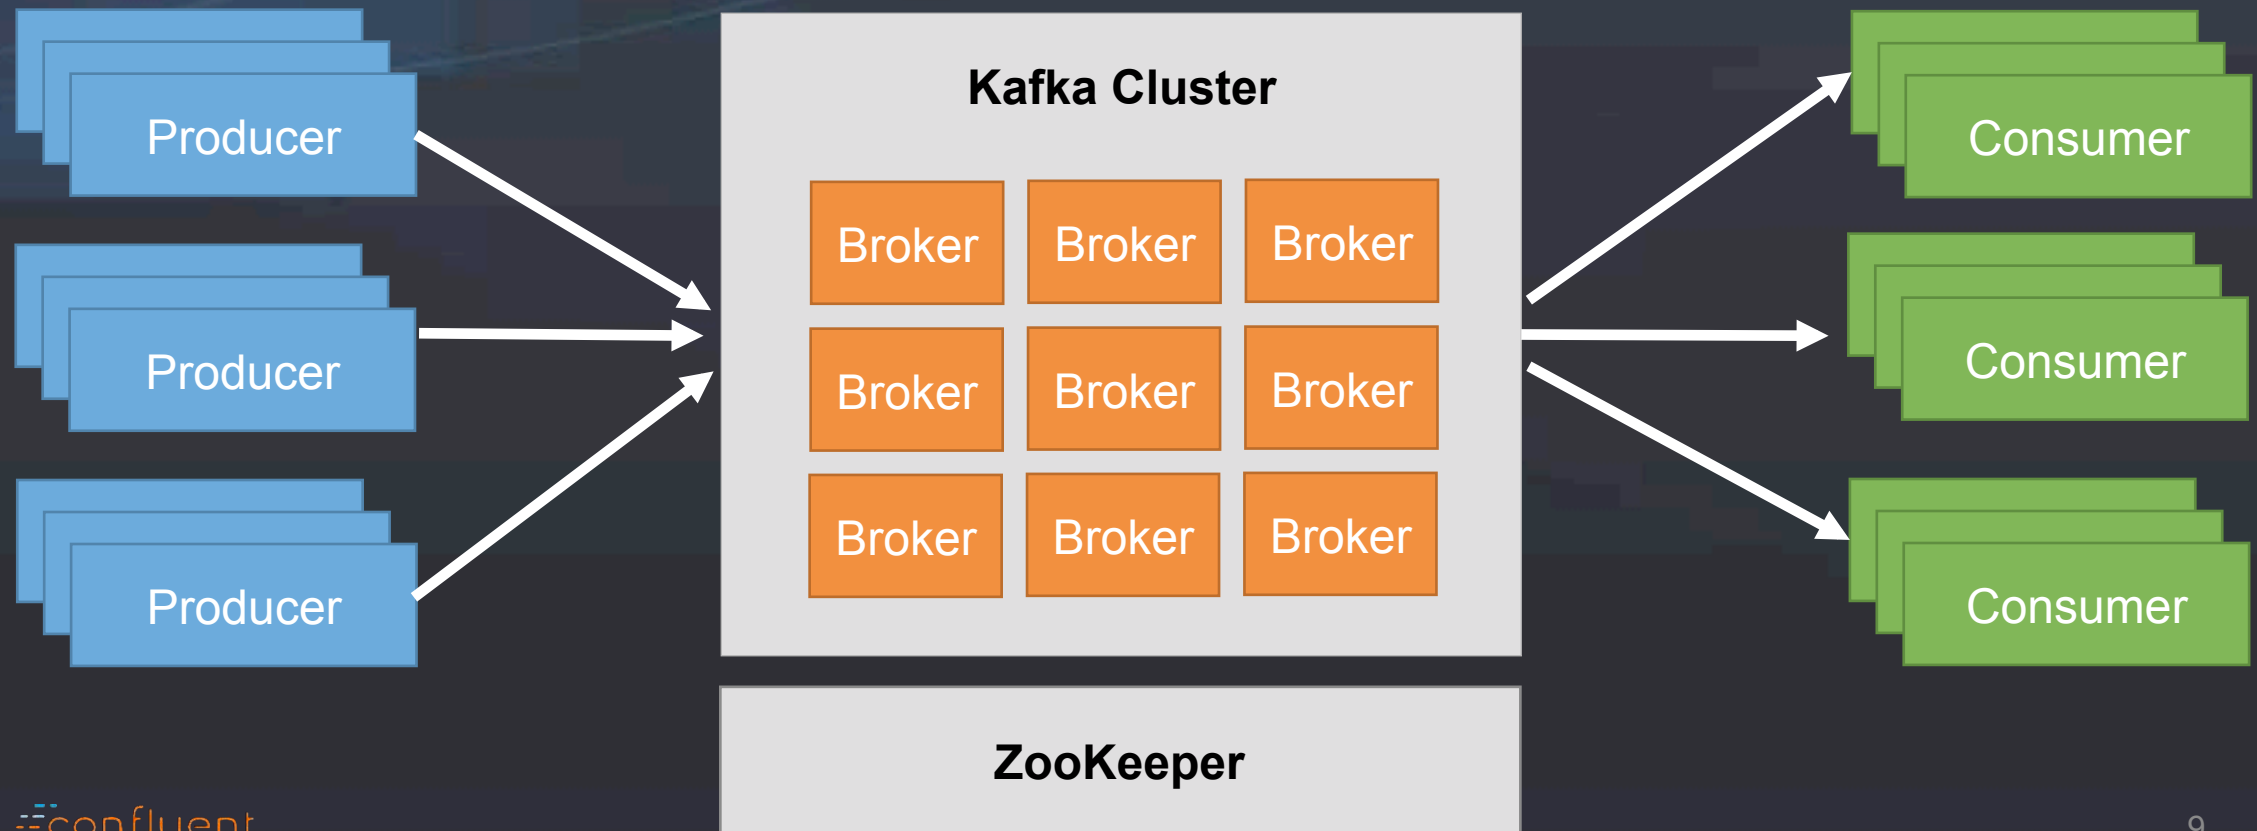


Apache Kafka is a **high-throughput** distributed messaging system.

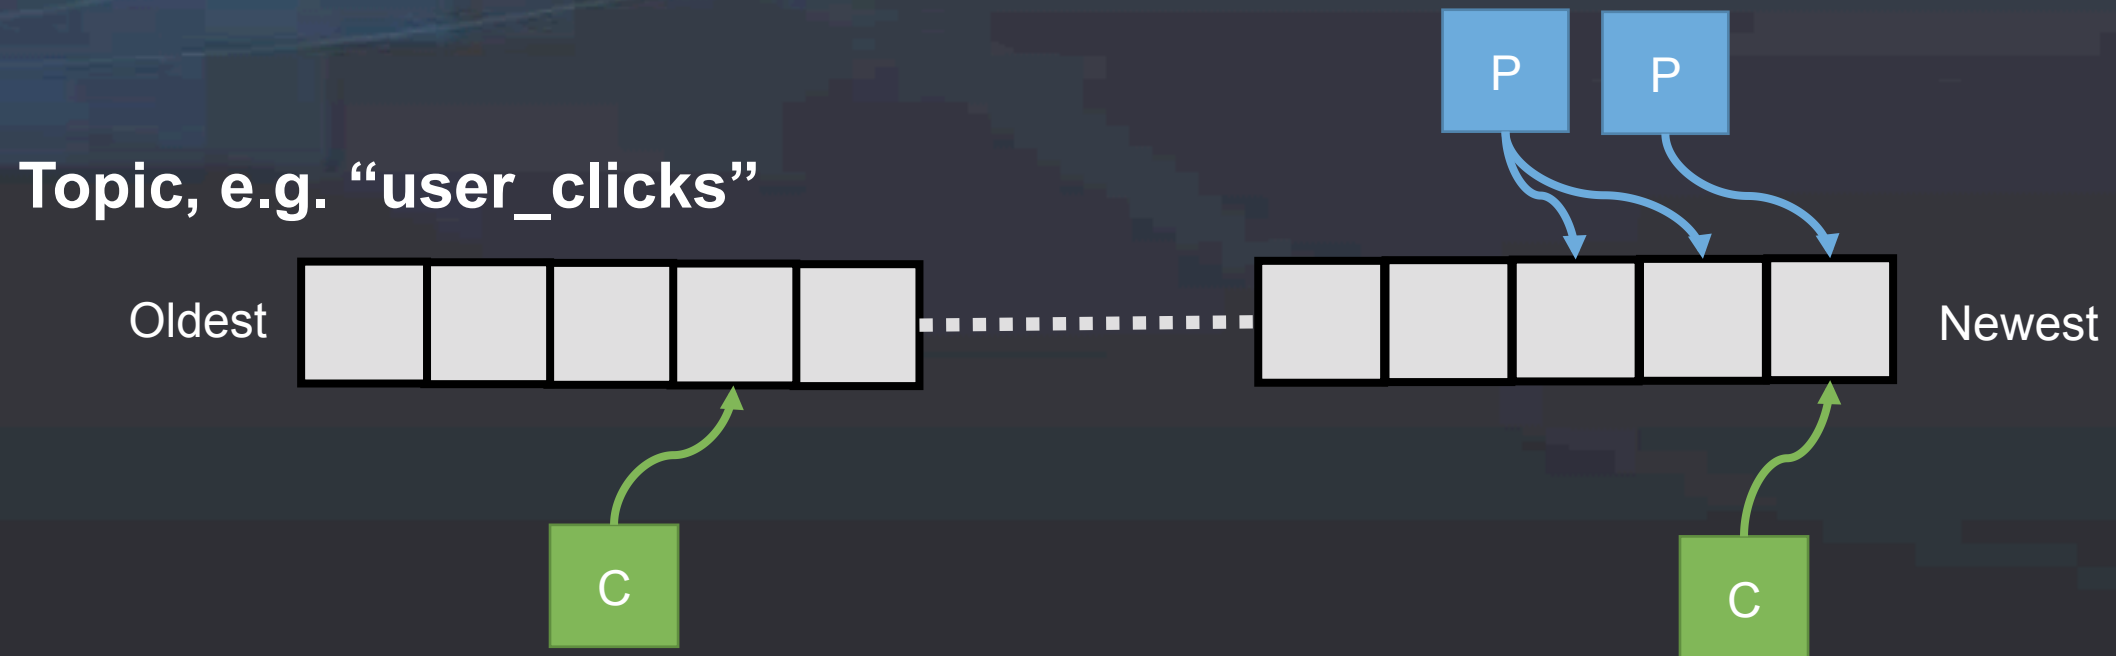
“1,100,000,000,000 msg/day, totaling 175+ TB/day” (LinkedIn)

= 3 billion messages since the beginning of this talk

Apache Kafka is a **publish-subscribe** messaging rethought as a **distributed commit log**.




Apache Kafka is a **publish-subscribe** messaging rethought as a **distributed commit log**.

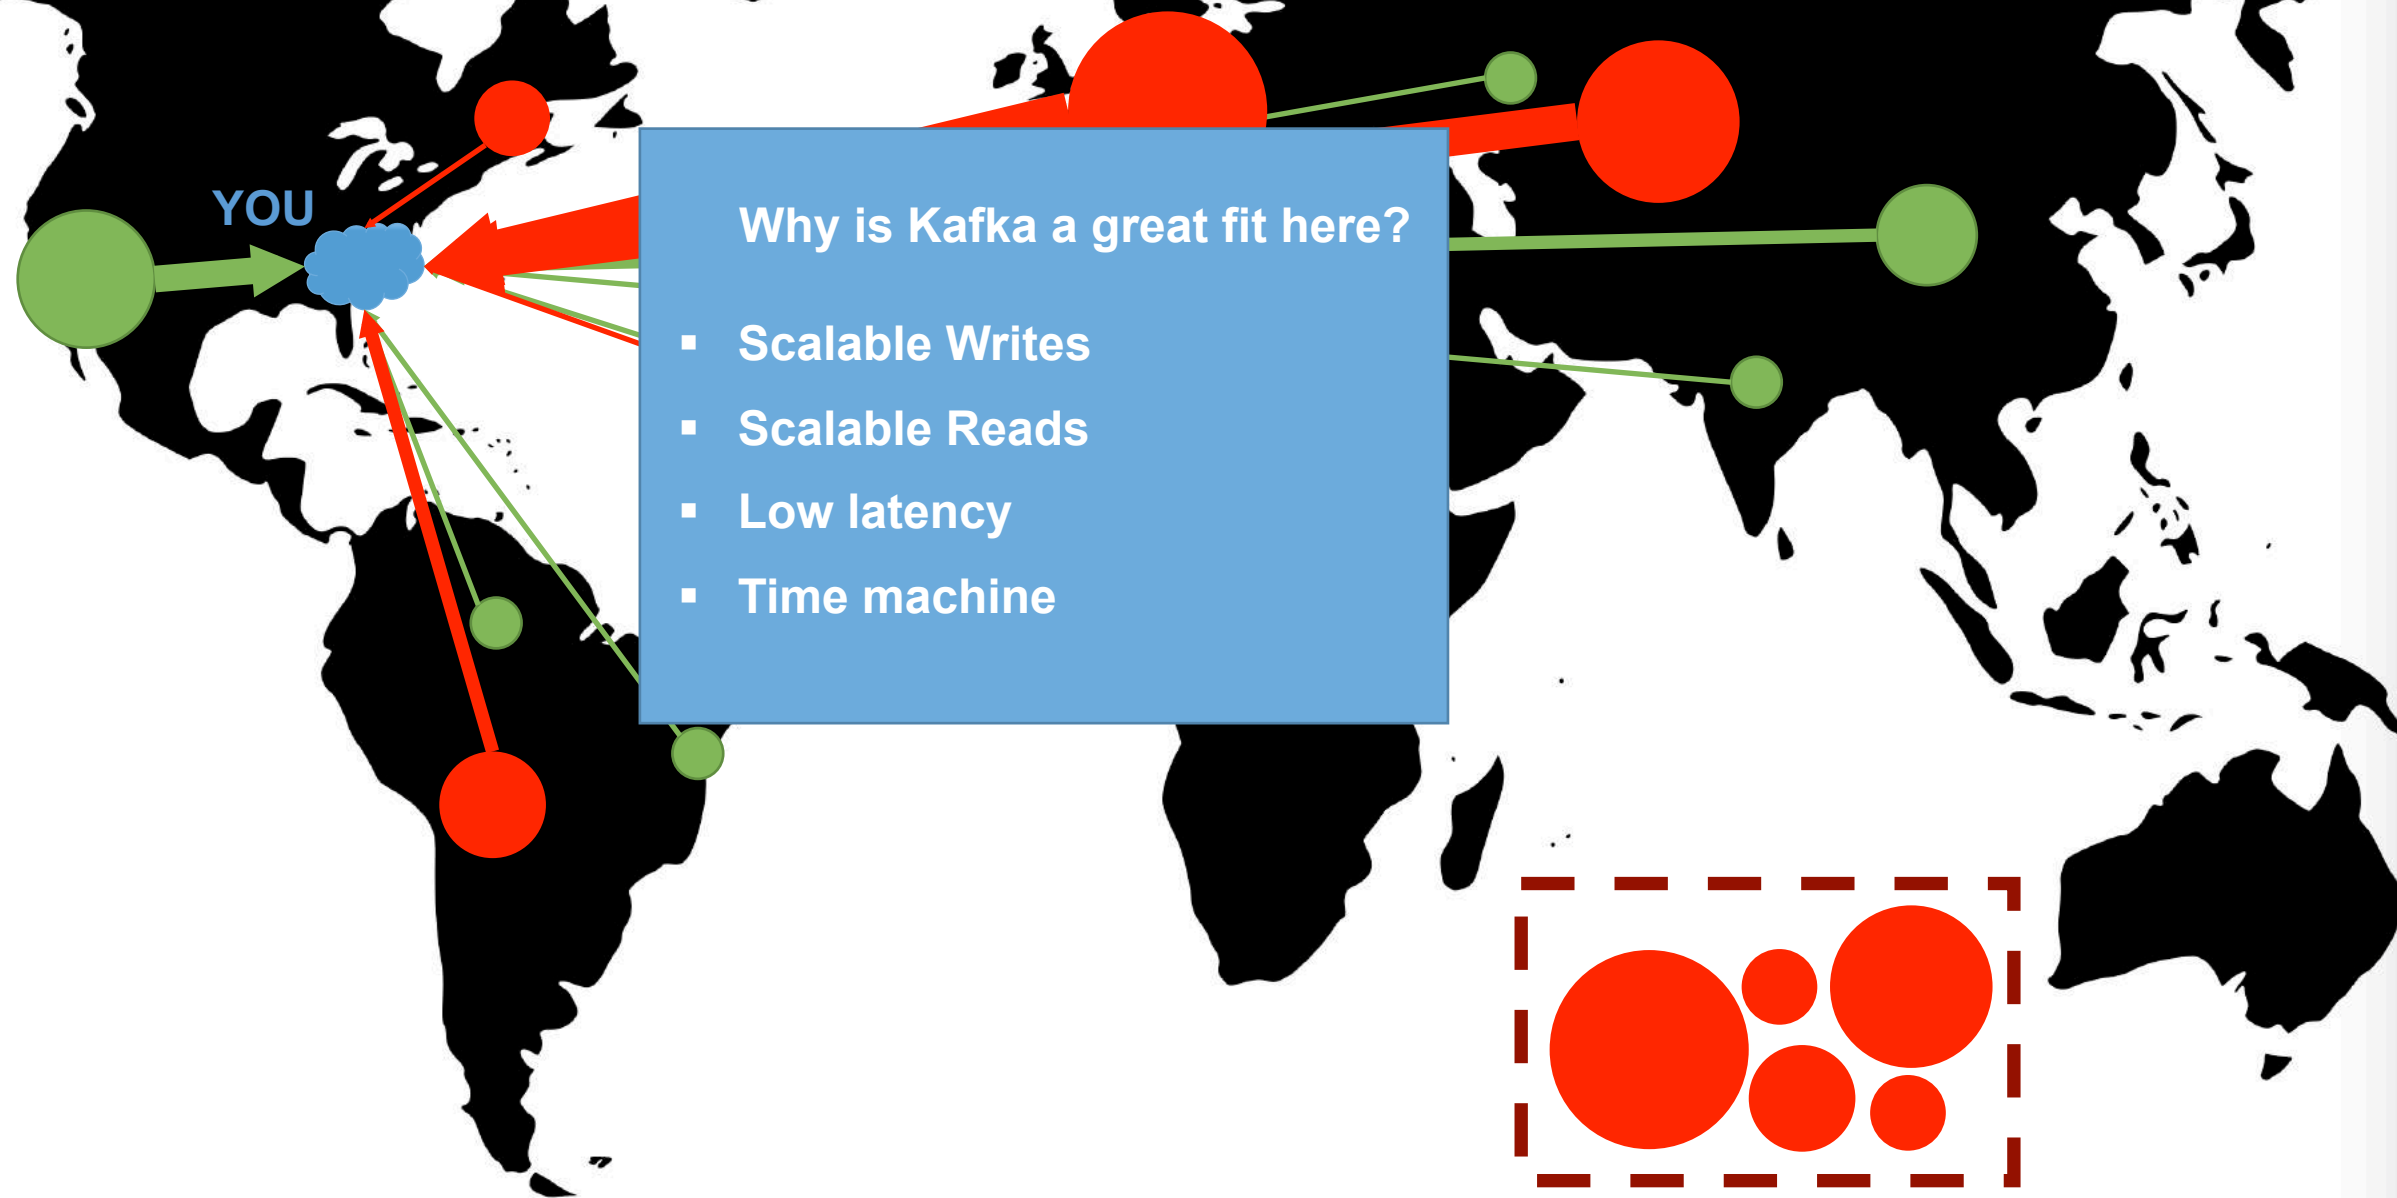


So where can Kafka help me?

Example, anyone?



Example: Protecting your infrastructure against DDoS attacks





Kafka powers many use cases

- User behavior, click stream analysis
- Infrastructure monitoring and security, e.g. DDoS detection
- Fraud detection
- Operational telemetry data from mobile devices and sensors
- Analyzing system and app logging data
- Internet of Things (IoT) applications
- ...and many more
 - Yours is missing? Let me know via michael@confluent.io !



Diverse and rapidly growing user base across many industries and verticals.

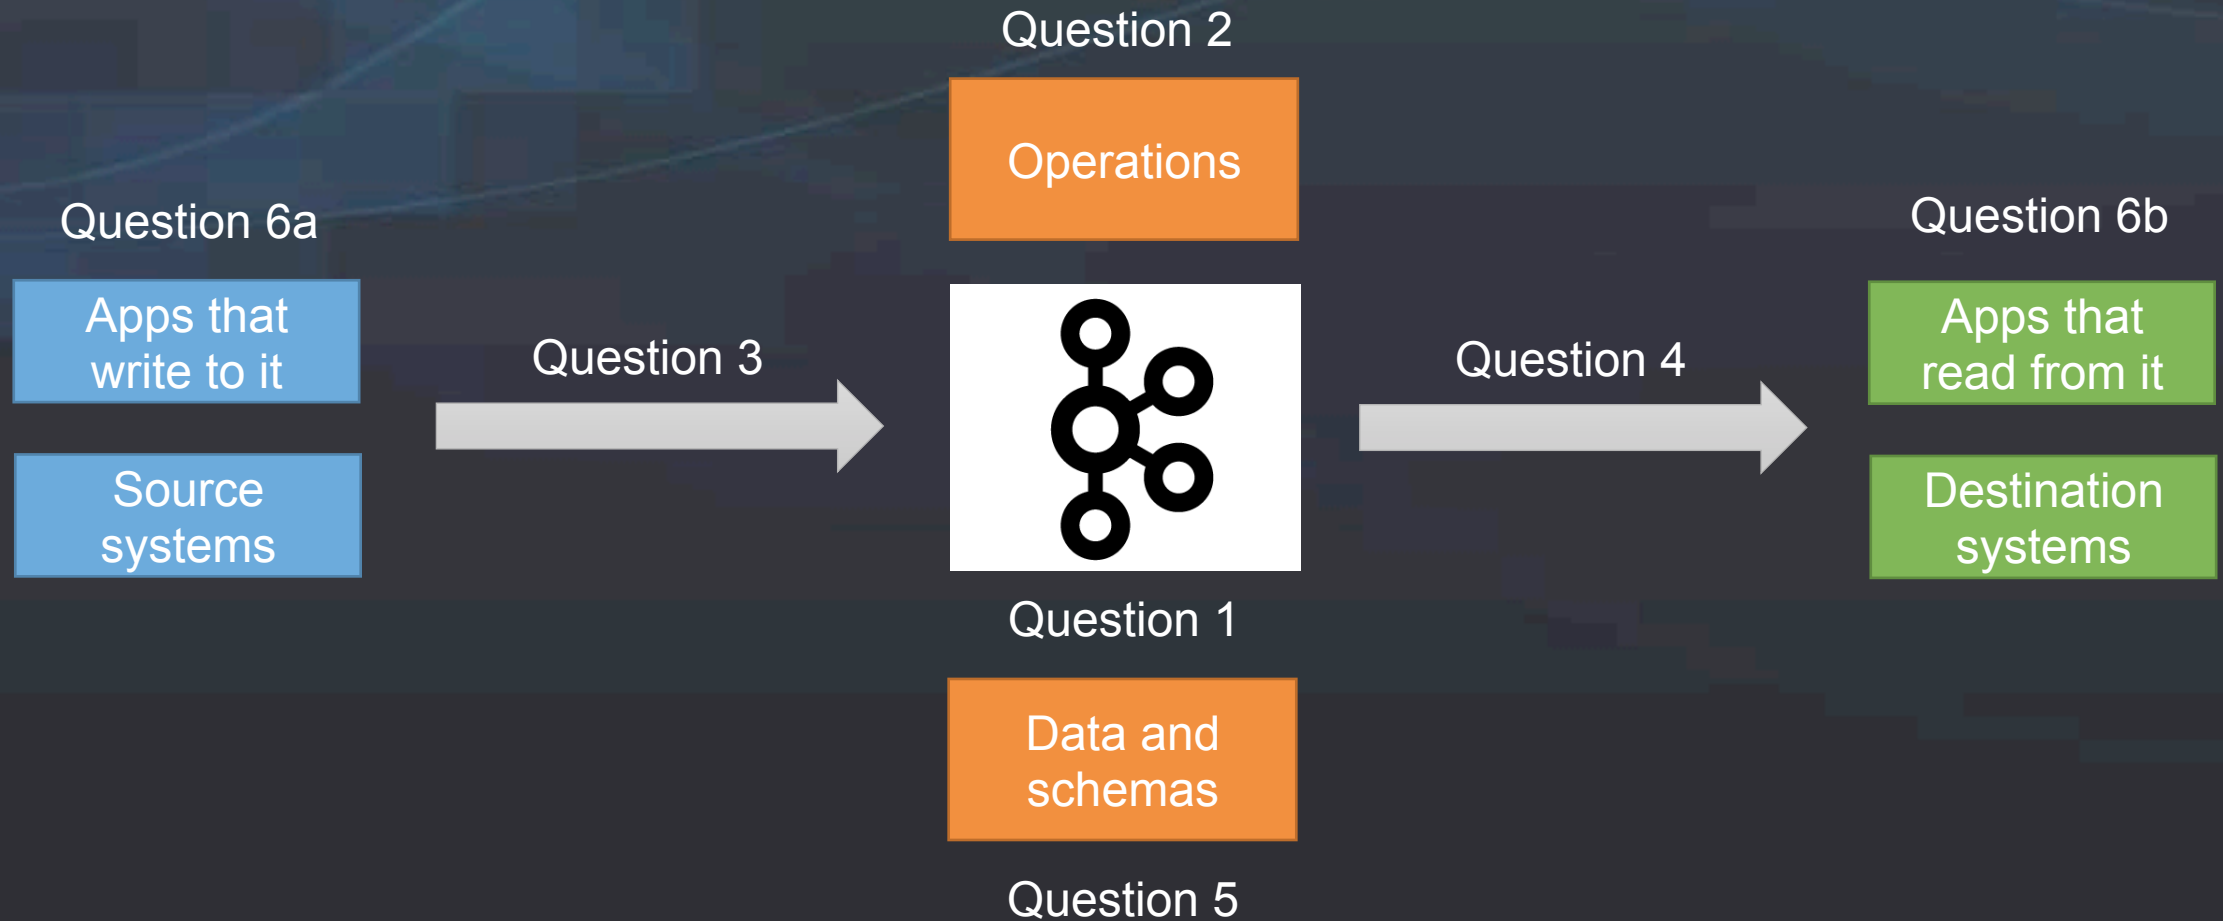


A typical Kafka architecture

Yes, we now begin to approach “production”

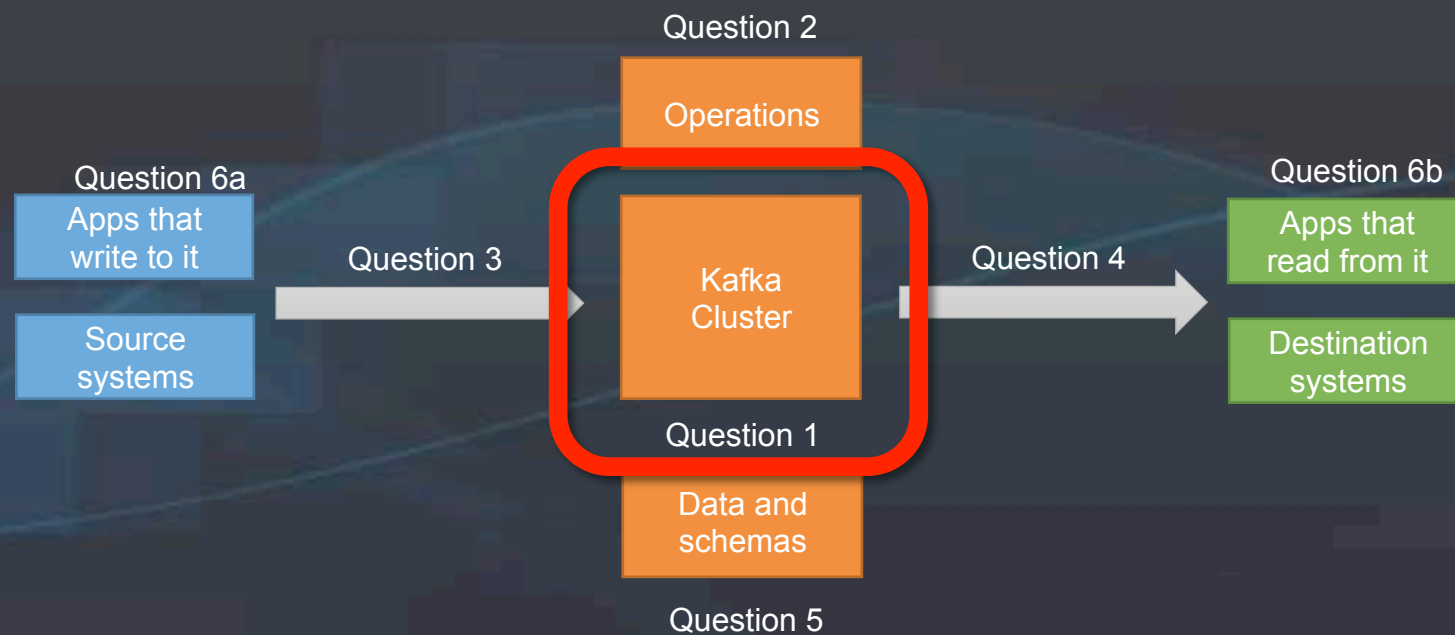


Typical architecture => typical questions



Wait a minute!





Kafka core

Question 1 or “What are the upcoming improvements to core Kafka?”

Kafka core: upcoming changes in 0.9.0

- Kafka 0.9.0 (formerly 0.8.3) expected in November 2015
- ZooKeeper now only required for Kafka **brokers**
 - ZK dependency removed from clients = producers and consumers
 - Benefits include less load on ZK, lower operational complexity, user apps don't require interacting with ZK anymore
- New, unified consumer Java API
 - We consolidated the previous “high-level” and “simple” consumer APIs
 - Old APIs still available and not deprecated (yet)

New consumer Java API in 0.9.0

```
import org.apache.kafka.clients.consumer.*;
```

```
Properties kafkaProps = new java.util.Properties();  
kafkaProps.put("group.id", "dump-to-console-app"); // plus any further Kafka settings
```

Configure

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<String, String>(kafkaProps);  
consumer.subscribe(java.util.Collections.singletonList("my-topic"));
```

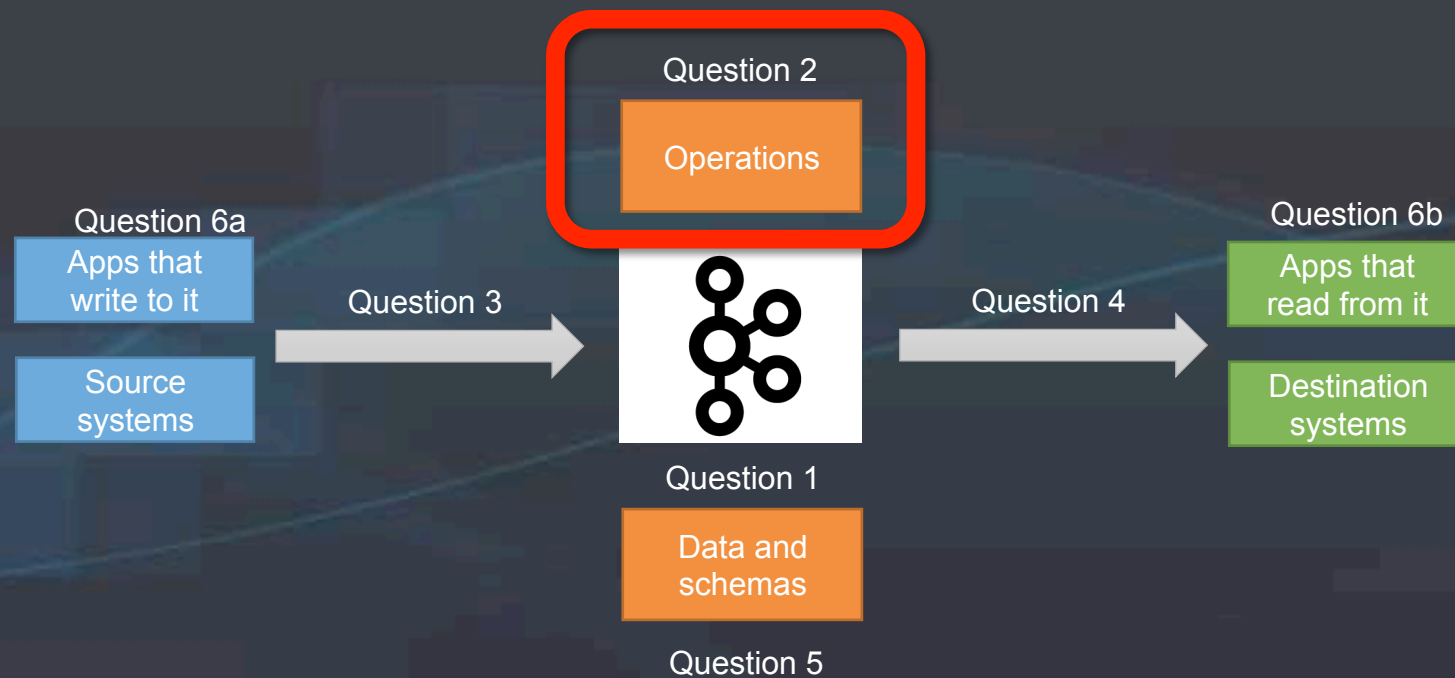
Subscribe

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records) {  
        System.out.printf("key = %s, value = %s\n", record.key(), record.value());  
    }  
}
```

Process

Kafka core: upcoming changes in 0.9.0

- Improved **data import/export** via Copycat
 - KIP-26: <https://cwiki.apache.org/confluence/pages/viewpage.action?pageId=58851767>
 - Will talk about this later
- Improved **security**: SSL support for encrypted data transfer
 - Yay, finally make your InfoSec team (a bit) happier!
 - <https://cwiki.apache.org/confluence/display/KAFKA/Deploying+SSL+for+Kafka>
- Improved **multi-tenancy**: quotas aka throttling for Ps and Cs
 - KIP-13: <https://cwiki.apache.org/confluence/display/KAFKA/KIP-13+-+Quotas>
 - Quotas are defined per broker, will slow down clients if needed
 - Reduces collateral damage caused by misbehaving apps/teams



Kafka operations

Question 2 or “How do I deploy, manage, monitor, etc. my Kafka clusters?”

Deploying Kafka

- Hardware recommendations, configuration settings, etc.
 - <http://docs.confluent.io/current/kafka/deployment.html>
 - <http://kafka.apache.org/documentation.html#hwandos>
- Deploying Kafka itself = DIY at the moment
- Packages for Debian and RHEL OS families available via Confluent Platform
 - <http://www.confluent.io/developer>
- Straight-forward to use orchestration tools like Puppet, Ansible
- Also: options for Docker, Mesos, YARN, ...

Managing Kafka: CLI tools

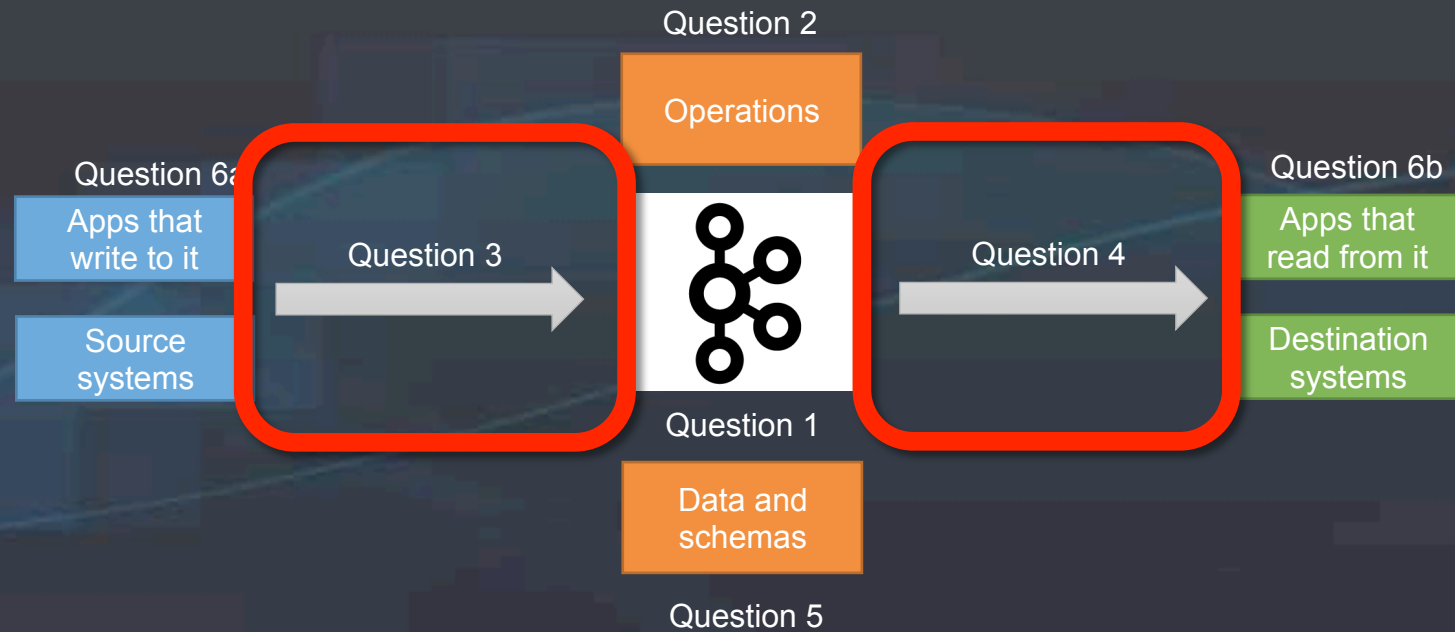
- Kafka includes a plethora of CLI tools
 - Managing topics, controlling replication, status of clients, ...
- Can be tricky to understand which tool to use, when, and how
- Helpful pointers:
 - <https://cwiki.apache.org/confluence/display/KAFKA/System+Tools>
 - <https://cwiki.apache.org/confluence/display/KAFKA/Replication+tools>
- KIP-4 will eventually add better management APIs

Monitoring Kafka: metrics

- How to monitor
 - Usual tools like Graphite, InfluxDB, statsd, Grafana, collectd, diamond
- What to monitor – some key metrics
 - Host metrics: CPU, memory, disk I/O and usage, network I/O
 - Kafka metrics: consumer lag, replication stats, message latency, Java GC
 - ZooKeeper metrics: latency of requests, #outstanding requests
- Kafka exposes many built-in metrics via JMX
 - Use e.g. jmxtrans to feed these metrics into Graphite, statsd, etc.

Monitoring Kafka: logging

- You can expect lots of logging data for larger Kafka clusters
- Centralized logging services help significantly
 - You have one already, right?
 - Elasticsearch/Kibana, Splunk, Loggly, ...
- Further information about operations and monitoring at:
 - <http://docs.confluent.io/current/kafka/monitoring.html>
 - <https://www.slideshare.net/miguno/apache-kafka-08-basic-training-verisign>



Kafka clients #1

Questions 3+4 or “How can my apps talk to Kafka?”

Recommended* Kafka clients as of today

Language	Name	Link
Java	<built-in>	http://kafka.apache.org/
C/C++	librdkafka	http://github.com/edenhill/librdkafka
Python	kafka-python	https://github.com/mumrah/kafka-python
Go	sarama	https://github.com/Shopify/sarama
Node	kafka-node	https://github.com/SOHU-Co/kafka-node/
Scala	reactive kafka	https://github.com/softwaremill/reactive-kafka
...

Polyglot Ready (tm)

Kafka clients: upcoming improvements

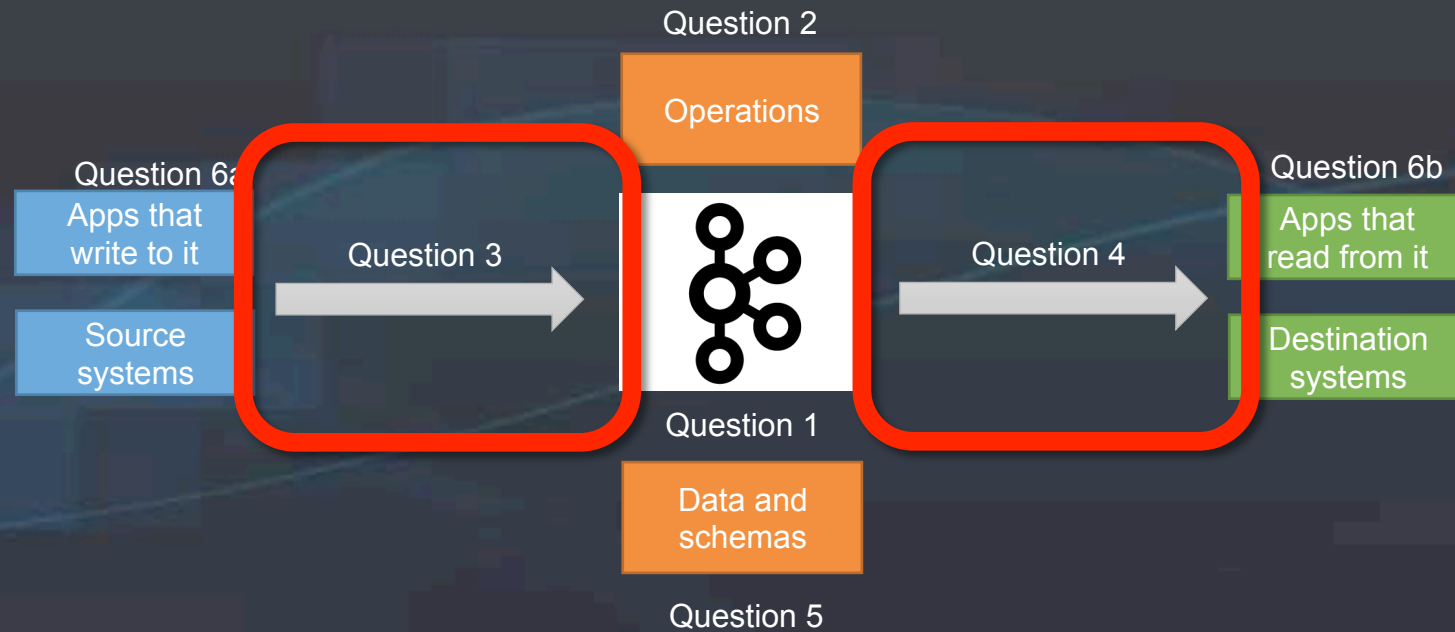
- Current problem: only **Java client** is officially supported
 - A lot of effort and duplication for client maintainers to be compatible with Kafka changes over time (e.g. protocol, ZK for offset management)
 - Wait time for users until “their” client library is ready for latest Kafka
- Idea: use **librdkafka** (C) as the basis for Kafka clients and provide bindings + idiomatic APIs per target language
- Benefits include:
 - Full protocol support, SSL, etc. needs to be implemented only once
 - All languages will benefit from the speed of the C implementation
 - Of course you are always free to pick your favorite client!

Confluent Kafka-REST

- Open source, included in Confluent Platform
<https://github.com/confluentinc/kafka-rest/>
- Alternative to native clients
- Supports reading and writing data, status info, Avro, etc.



```
# Get a list of topics
$ curl "http://rest-proxy:8082/topics"

[ { "name":"userProfiles",    "num_partitions": 3 },
  { "name":"locationUpdates", "num_partitions": 1 } ]
```

Kafka clients #2

Questions 3+4 or “How can my systems talk to Kafka?”



Data import/export: status quo

- Until now this has been **your** problem to solve
 - Only few tools available, e.g. LinkedIn Camus for Kafka → HDFS export
 - Typically a DIY solution using the aforementioned client libs
- Kafka 0.9.0 will introduce **Copypcat**

Copypcat is the **I/O redirection** in your Unix pipelines.
Use it to get your data into and out of Kafka.

```
$ cat < in.txt | grep "apache" | tr a-z A-Z > out.txt
```



this



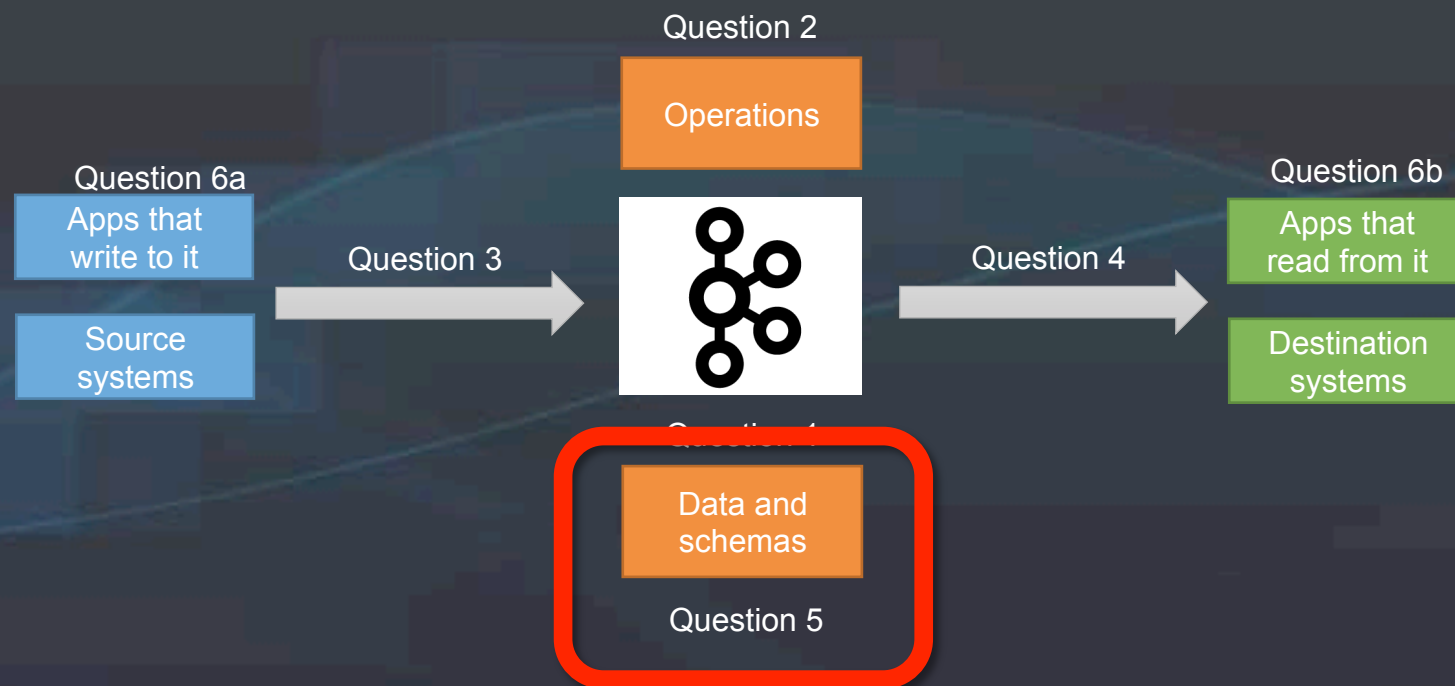
this

Data import/export via Copycat

- **Copycat** is included in upcoming Kafka 0.9.0



- Federated Copycat “connector” development for e.g. HDFS, JDBC
- Light-weight, scales from simple testing and one-off jobs to large-scale production scenarios serving an entire organization
- Process management agnostic, hence flexible deployment options
 - Examples: Standalone, YARN, Mesos, or your own (e.g. Puppet w/ supervisord)



Data and schemas

Question 5 or “Je te comprends pas”

ALL I SAID WAS



**LET'S AGREE ON A DATA
FORMAT**

memegenerator.net

BONJOUR!



WTF?



Data and schemas

- Agree on **contracts for data** just like you do for, say, APIs
 - Producers and consumers of data must understand each other
 - Free-for-alls typically degenerate quickly into team deathmatches
 - Benefit from clear contract, schema evolution, type safety, etc.
- Organizational problem rather than technical
 - Hilarious /facepalm moments
 - Been there, done that 😊
- Take a look at **Apache Avro**, Thrift, Protocol Buffers
 - Cf. Avro homepage, <https://github.com/miguno/avro-hadoop-starter>

“Alternative” to schemas



Example: Avro schema for tweets

```
{  
  "type": "record",  
  "name": "Tweet",  
  "namespace": "io.confluent.avro",  
  "fields": [  
    {  
      "name": "username",  
      "type": "string",  
      "doc": "Name of the user account on Twitter.com"  
    },  
    {  
      "name": "tweet",  
      "type": "string",  
      "doc": "The content of the user's Twitter message"  
    },  
    {  
      "name": "timestamp",  
      "type": "long",  
      "doc": "Unix epoch time in seconds"  
    }  
  ],  
  "doc": "A basic schema for storing Twitter messages"  
}
```

username

text

timestamp

<data> = <definition>

<data> = <definition>

<data> = <definition>




<data> = <definition>

<data> = <definition>

“UserProfile” = <definition>


“Tweet” = <definition>

“Alert” = <definition>



Schema registry

- Stores and retrieves your schemas
- Cornerstone for building resilient data pipelines
- Viable registry implementation missing until recently
 - AVRO-1124 (2012-2014)
 - So everyone started to roll their own schema registry
 - Again: been there, done that 😊
- There must be a better way, right?

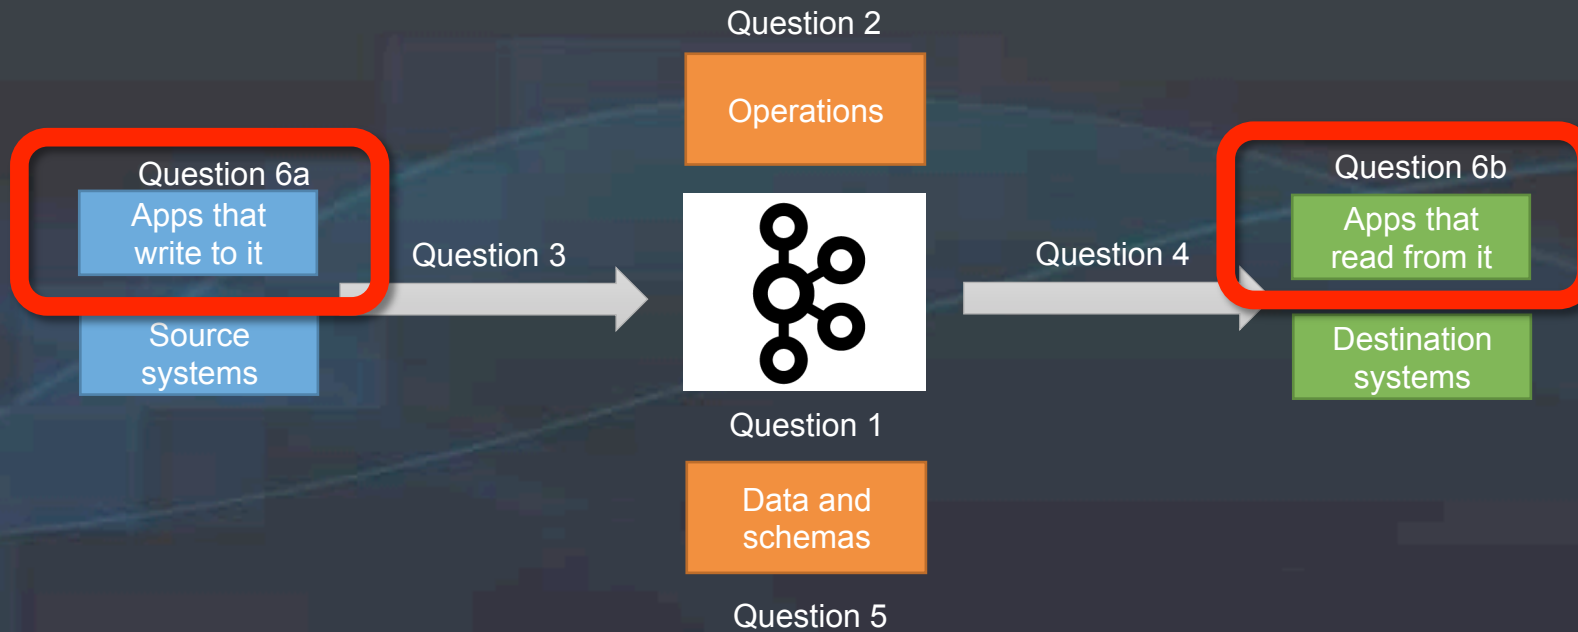


Confluent Schema Registry

- Open source, included in Confluent Platform
<https://github.com/confluentinc/schema-registry/>
- REST API to store and retrieve schemas etc.

```
# List all schema versions registered for topic "foo"  
$ curl -X GET -i http://registry:8081/subjects/foo/versions
```

- Integrates with Kafka clients, Kafka REST, Camus, ...
- Can enforce policies for your data, e.g. backwards compatibility
- Still not convinced you need a schema registry?
 - <http://www.confluent.io/blog/schema-registry-kafka-stream-processing-yes-virginia-you-really-need-one>



Stream processing

Question 6 or “How do I actually process my data in Kafka?”

Stream processing

- Currently three main options
 - Storm: arguably powering the majority of production deployments
 - Spark Streaming: runner-up, but gaining momentum due to “main” Spark
 - DIY: write your own using Kafka client libs, typically with a narrower focus

Some people, when confronted with a problem to process data in Kafka, think “I know, I’ll use [Storm | Spark | ...].”

Now they have *two* problems.

Stream processing

Four!

- Currently ~~three~~ main options
 - Storm: arguably powering the majority of production deployments
 - Spark Streaming: runner-up, but gaining momentum due to “main” Spark
 - DIY: write your own using Kafka client libs, typically with a narrower focus
- Kafka 0.9.0 will introduce **Kafka Streams**

Kafka Streams is the **commands** in your Unix pipelines.
Use it to transform data stored in Kafka.

```
$ cat < in.txt | grep "apache" | tr a-z A-Z > out.txt
```



this



this

Kafka Streams

- Kafka Streams included in Kafka 0.9.0
 - KIP-28: <https://cwiki.apache.org/confluence/display/KAFKA/KIP-28+-+Add+a+processor+client>
- No need to run another framework like Storm alongside Kafka
 - No need for separate infrastructure and trainings either
- **Library** rather than framework
 - Won't dictate your deployment, configuration management, packaging, ...
 - Use it like you'd use Apache Commons, Google Guava, etc.
 - Easier to integrate with your existing apps and services
- 100% compatible with Kafka by definition

Kafka Streams

- Initial version will support
 - Low-level API as well as higher-level API for Java 7+
 - Operations such as join/filter/map/...
 - Windowing
 - Proper time modeling, e.g. event time vs. processing time
 - Local state management with persistence and replication
 - Schema and Avro support
- And more to come – details will be shared over the next weeks!


Example of higher-level API (much nicer with Java 8 and lambdas)

```
KStream<String, Integer> stream2 =  
    stream1.map(new KeyValueMapper<String, String, KeyValue<String, Integer>>() {  
        @Override  
        public KeyValue<String, Integer> apply(String key, String value) {  
            return new KeyValue<>(key, new Integer(value));  
        }  
    }).filter(new Predicate<String, Integer>() {  
        @Override  
        public boolean apply(String key, Integer value) {  
            return true;  
        }  
    });
```

map()

filter()

Phew, we made it!



Copycat

Kafka Streams

```
$ cat < in.txt | grep "apache" | tr a-z A-Z > out.txt
```

Kafka

Want to contribute to Kafka and open source?

Join the Kafka community
<http://kafka.apache.org/>

...in a great team with the creators of Kafka
and also getting paid for it?

Confluent is hiring 😊
<http://confluent.io/>

Questions, comments? Tweet with **#ApacheBigData** and /cc to **@ConfluentInc**