

# 深入了解Redis

宋传胜

[Chuansheng.song@langtaojin.com](mailto:Chuansheng.song@langtaojin.com)



# 简单介绍

- 文本协议 memcached类似
- KEY 可打印字符
- VALUE支持的类型
  - STRINGS
  - LIST
  - SET
  - SORTED SET
  - HASH
- 高性能 (100k+ SET/80k+ GET)/s
- 序列化支持
- 主从同步支持
- 客户端自己实现sharding

# 基本数据结构

- RedisObject (Redis.h)
  - key/value对象的基础

```
typedef struct redisObject {  
    unsigned type:4;  
    unsigned notused:2;    /* Not used */  
    unsigned encoding:4;  
    unsigned lru:22;    /* lru time (relative to server.lruclock) */  
    int refcount;  
    void *ptr;  
    /* VM fields are only allocated if VM is active, otherwise the  
    * object allocation function will just allocate  
    * sizeof(redisObjct) minus sizeof(redisObjectVM), so using  
    * Redis without VM active will not have any overhead. */  
} robj;
```



# KEY

- 基本命令

- get/set/exists  $O(1)$
- setnx/randomkey/rename/renamenx/dbsize/type
- keys pattern
- del 一次删除多个key
- expire key seconds 指定过期时间
- ttl key 剩余过期时间
- select db
- move key newdb
- flush db/flushall

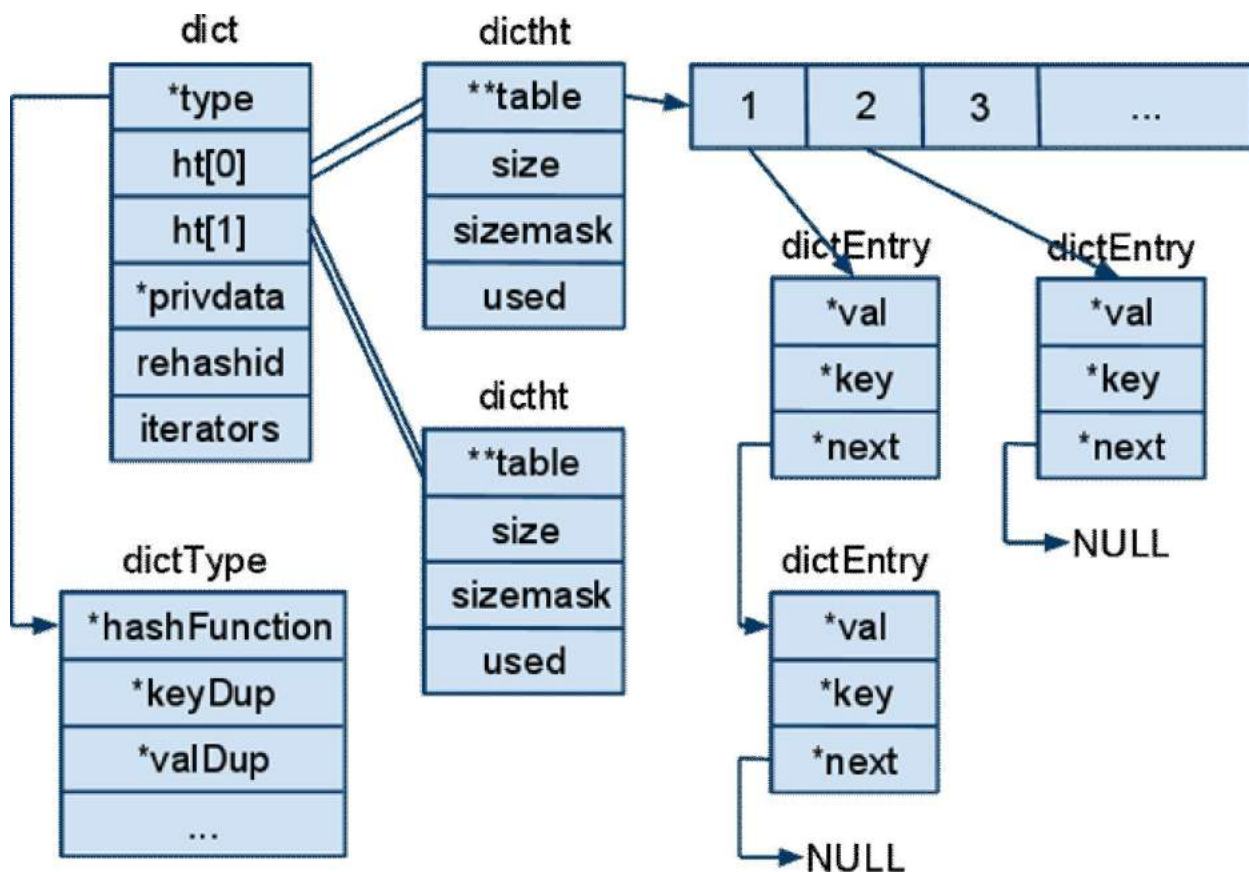


# EXPIRE AND LRU CACHE

- **Volatile Key (Db.c)**
  - When the key is set to a new value using the [SET](#) command, or when a key is destroyed via [DEL](#), the timeout is removed from the key.
- **Enhanced Lazy Expiration algorithm**
  - Redis does not constantly monitor keys that are going to be expired. Keys are expired simply when some client tries to access a key, and the key is found to be timed out.
- **Work as LRU cache (memcached)**
  - Maxmemory/maxmemory-policy
  - When memory limit was already reached, server will remove some old data deleting a *volatile* key, even if the key is still far from expiring automatically.
  - Random get keys, delete by lru rules.

# 基本数据结构

- Hashtable(Dict.c)





# 基本Value数据结构

- STRING (sds.c)

```
struct sdshdr {  
    int len;  
    int free;  
    char buf[];  
};
```

- set/setnx/get/getset/mget/mset/msetnx
  - nx – not exists
  - m – multiple
- incr/decr/incrby/decrby
- getrange/append

# 基本Value数据结构

- LIST (T\_list.c)
  - REDIS\_LIST 类型, 如果其 entry 小于配置值: list-max-ziplist-entries 或 value字符串的长度小于 list-max-ziplist-value, 使用ziplist数据结构, 否则使用标准的Doubly linked list
  - (l/r)push/(l/r)pop/lLen  $O(1)$
  - b(l/r)pop支持阻塞等待, 避免了轮循
  - lrange/ltrim/lrem/lset/rpoplpush



# LIST

- Ziplist (Ziplist.c)
  - $O(\text{mem\_size})$  add/delete
  - list-max-ziplist-entries (default: 1024)
  - list-max-ziplist-value (default: 32)




# LIST

- Ziplist (continue)
  - zlentry (unencode structure)

```
typedef struct zlentry {  
    unsigned int prevrawlensize, prevrawlen;  
    unsigned int lensize, len;  
    unsigned int headersize;  
    unsigned char encoding;  
    unsigned char *p;  
} zlentry;
```

# LIST

- Doubly linked list (Adlist.c)



# 基本Value数据结构

- SET (T\_set.c) – hashtable + Intset
  - String的无序集合
  - sadd/srem/sismember/scard  $O(1)$
  - spop/srandmember
  - smove
  - smembers
  - sinter(store)  $O(C)$
  - sunion(store)/sdiff(store)  $O(N)$

# INTSET

- INTSET (Intset.c)
  - 都是整型时数据结构退化成排序的intset
  - Good fit for size up to 20-50K
  - set-max-intset-entries (default: 4096)
  - $O(n)$  search
  - $O(\log(n) + \text{mem\_size})$  add/del

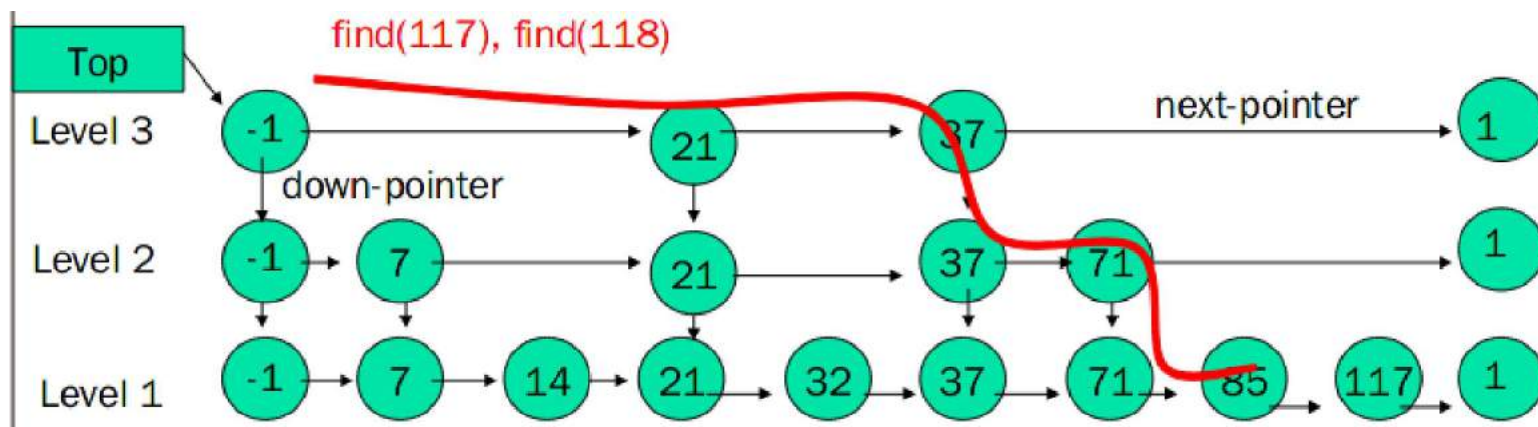
```
typedef struct intset {  
    uint32_t encoding;  
    uint32_t length;  
    int8_t contents[];  
} intset;
```

# 基本Value数据结构

- SORTED SET(T\_zset.c) hashtable + skip list
  - 按照key的score排序
  - zadd/zrem/zrank/zrevrank  $O(\log(n))$
  - zcard  $O(1)$
  - zincrby 修改score
  - zrange/zrevrange/zrangebyscore/zscore  $O(\log(N)+M)$
  - zremrangebyrank/zremrangebyscore  $O(\log(N)+M)$

# SKIP LIST

- Skip List (T\_zset.c)
  - skip list → redblack tree → AVL tree (more balanced)
  - lockfree/concurrency






# 基本Value数据结构

- HASH (T\_hash.c) zipmap + hashtable
  - 如果其 entry 小于配置值: hash-max-zipmap-entries 或 value字符串的长度小于 hash-max-zipmap-value , 使用zipmap数据结构
  - Value只能是string类型
  - hset/hget/hexists  $O(1)$
  - hmset/hexists/hincrby/hdel/hlen
  - hmget  $O(N)$
  - hkeys/hvals/hgetall

# ZIPMAP

- zipmap
  - $O(n)$  search  $O(\text{mem\_size})$  add/delete



# 内存管理

- `zmalloc.c`
  - 简单内存管理
  - 支持`tcmalloc` `USE_TCMALLOC`
  - `zmalloc(size_t size) → malloc(size+PREFIX_SIZE)`
  - `fragmentation_ratio =`  
`zmalloc_get_rss/zmalloc_used_memory`



# PUBSUB

- SUBSCRIBE/UNSUBSCRIBE/PSUBSCRIBE/PUNSUBSCRIBE
- PUBLISH
- 连接断开后信息丢失

# TRANSACTION


- MULTI （不支持CAS）
  - MULTI
  - SET foo 0
  - INCR foo
  - EXEC
- WATCH （Check and Set ）
  - WATCH theKey
  - v = GET theKey
  - MULTI
  - SET theKey v+1
  - EXEC
- UNWATCH/DISCARD

# 提升网络效率、减少内存占用

- MGET/MSET
- PIPELINING
  - `$ (echo -en "PING\r\nPING\r\nPING\r\n"; sleep 1) | nc localhost 6379`  
+PONG  
+PONG  
+PONG
- 压缩数据，触发ziplist/zipmap/intset
- 使用 GETBIT/SETBIT/GETRANGE/SETRANGE 压缩多个属性
- 使用HASH
- 使用32位redis + sharding

# Client Side Sharding

- Consistent Hashing
  - Redis Sharding at Craigslist



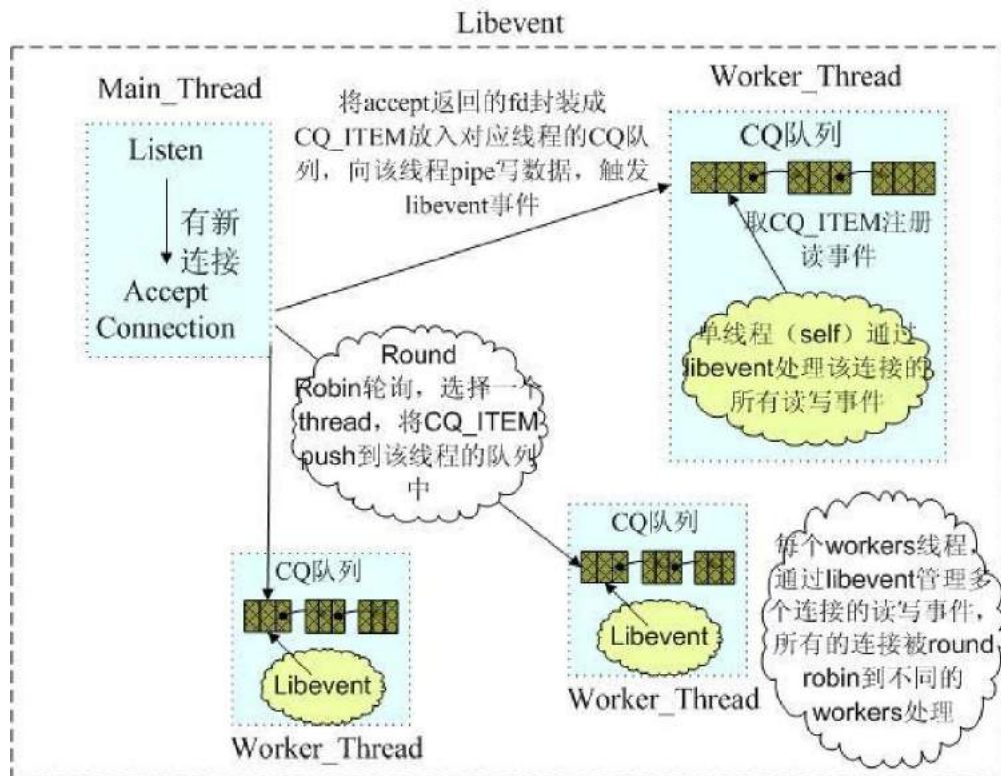


# 网络层

- Anet.c
- 事件驱动 ae.c →  
ae\_epoll.c/ae\_kqueue.c/ae\_select.c
- 非阻塞单线程
- 支持timer
- 在VM场景下有可能多线程

# 网络层对比

- memcached/varnish/scribed的网络IO模型
  - 多线程，非阻塞IO复用



# 网络层对比

- lighttpd/nginx的网络IO模型
  - 多进程，单线程非阻塞IO
- apache的MPM模型
  - 多进程prefork，多线程

# PERSISTENCE


- SNAPSHOT (Rdb.c)
- 可以关闭
- BGSAVE or save xseconds ychanges
- fork
  - parent - disable rehash, periodic check child on serverCron (100ms)
  - child - copy-on-write, save whole db, exit
- Need to turn on the overcommit\_memory setting if you want to deal with a dataset more than 1/2 of RAM, which contradicts our habit to battle against OOM killer as a database administrator.

# Append Only File

- Append to log file on every change (Aof.c)
- fsync() policy
  - always/os decide/every second
  - Always with battery backed raid controllers
  - Every second by default (innodb\_flush\_log\_at\_trx\_commit=2 in mysql)
- Compact aof file
  - BGREWRITEAOF/REWRITEAOF
- Fork
  - child – write new aof in temporary file
  - parent – write change in both old aof file and memory buffer, append memory buffer to temporary file when child done, rename temporary to aof file.
- redis-check-aof --fix <filename>

# Replication

- 支持多级同步 (Replication.c)



# Replication

- 支持AUTH
- 起动过程
  - slave
    - SYNC
    - Wait
  - master
    - Issue a BGSAVE, or wait if BGSAVE in progress



# Virtual Memory

- Redis VM is now deprecated. Redis 2.4 will be the latest Redis version featuring Virtual Memory!
- Why not OS VM
  - Varnish [What's wrong with 1975 programming?](#)
  - Redis [What's wrong with 2006 programming?](#)
    - Blocking when page fault
    - 4k pages granularity
    - Optimal disk persistent format, 10:1
- Keys can't be swapped out
- vm is read only when BGSAVE/BGREWRITEAOF in progress
- vm-max-memory/vm-pages/vm-page-size

# Virtual Memory

- Blocking vm
- Threaded vm
  - Before command is executed, check the value
  - Block client and load, then signal by unix pipe
  - Continue
- Recommended filesystem – ext3 or any other file system with good support for *sparse files*
- redis-stat vmstat
- InnoDB innodb\_buffer\_pool\_size – swap and durability



# Future

- Diskstore (diskstore.c dscache.c)
  - All data on disk, B+tree in future
  - Swaps data into memory when needed
- Leveldb/InnoDB Change Buffer
  - LSM tree

# 其他相关的代码

- quicksort (pqsort.c)
- lzf compression (lzf\_c.c lzf\_d.c)
- Sha1 (Sha1.c)
- Syncio.c
- SORT (Sort.c)

# 线上Redis集群情况

- 16台服务器，兼做后端服务器
- 45个Redis Instance进程，三组服务
- 总内存 77G VIRT, 71G RSS
- 使用了主从复制、snapshot持久化，VM支持，客户端sharding、pipelining、hash优化

# Redis应用

- Resque – message queue of github (celery)
- Nginx HttpRedis module
- Redweb - Web administrative and query UI
- A fast, fuzzy, full-text index using Redis
- Redis-backed BloomFilter(s) in Ruby
- [Soulmate : Auto Complete with Redis](#)
- [Rate limiting with Redis](#)
- Redis Sharding
- Locking with setnx
- A Collection Of Redis Use Cases
- Radishapp Redis monitoring
- openredis.com hosting service

# Contribute to Redis

- Fork redis on github
  - <https://github.com/antirez/redis>
- Choose branch and commit your changes
- Create a pull request
- Wait for approval
  - <https://github.com/antirez/redis/pulls>
- My contribute for read only slave
  - <https://github.com/antirez/redis/pull/47>





# Resources

- DOCUMENTATION
  - <http://redis.io/documentation>
  - <http://antirez.com/>
  - <http://groups.google.com/group/redis-db>
- 本文参考的Redis资料集合
  - <http://www.delicious.com/fakechris/redis>
  - <http://www.delicious.com/tag/redis>